Philosophia (2008) 36:141–150
DOI 10.1007/s11406-007-9094-4

A Dilemma for Rule-Consequentialism

Jussi Suikkanen

Received: 30 April 2007 / Revised: 23 July 2007 / Accepted: 7 August 2007 /
Published online: 3 October 2007

Abstract Rule-consequentialists tend to argue for their normative theory by claiming that their view matches our moral convictions just as well as a pluralist set of Rossian duties. As an additional advantage, rule-consequentialism offers a unifying justification for these duties. I challenge the first part of the rule-consequentialist argument and show that Rossian duties match our moral convictions better than the rule-consequentialist principles. I ask the rule-consequentialists a simple question. In the case that circumstances change, is the wrongness of acts determined by the ideal principles for the earlier circumstances or by the ideal ones for the new circumstances? I argue that whichever answer the rule-consequentialists give the view leads to normative conclusions that conflict with our moral intuitions. Because some set of Rossian duties can avoid similar problems, rule-consequentialism fails in the reflective equilibrium test advocated by the rule-consequentialists.

Keywords Rule-consequentialism · Rossian pluralism · Reflective equilibrium test

Introduction

It is likely that a majority of ethicists agree that a crucial test for assessing moral theories is to ask to what extent their prescriptions and permissions cohere with our carefully reflected moral convictions (see, for instance, Daniels 1979; Hooker 1996, 2000a, pp. 9–16; Rawls 1971 pp. 19–21 and pp. 46–51; Ross 1930, p. 40; and Williams 1985, p. 94). According to this narrow reflective equilibrium test, if of two normative moral theories one coheres better with our moral convictions than the other, we should not accept the second one as the correct normative theory.[1]

[1] This test of narrow reflective equilibrium is conditional on the fact that the theory which fares worse in it does not cohere considerably better with our non-moral beliefs, i.e., on the fact that this theory does not fare significantly better in the wide reflective equilibrium test.

J. Suikkanen (✉)
Philosophy Department, University of Reading,
Reading RG6 6AA, UK
e-mail: jussi.suikkanen@helsinki.fi

I will argue that, if we are committed to this test, we should not accept rule-consequentialism. This is because rule-consequentialism faces a dilemma when we pose a simple question about which principles determine the wrongness of acts in different circumstances. Whichever answer the rule-consequentialists give to this question, there are bound to be cases where rule-consequentialism leads to certain counterintuitive moral prescriptions and permissions.[2] After showing this, I will argue that another moral theory, some version of Rossian pluralism, can avoid these awkward consequences. Given the comparative reflective equilibrium test, we should therefore reject rule-consequentialism.

My criticism of rule-consequentialism will be based on premises to which many rule-consequentialists are themselves committed. They often use the reflective equilibrium test described above to argue for their view and against Rossian pluralism (Hooker 1996 and 2000a, pp. 104–5). The argument is that the rule-consequentialist set of moral principles matches our moral convictions just as well as the Rossian set of *prima facie* duties. However, rule-consequentialism, in addition, provides a unified justification for these duties. Because this is an important virtue of a moral theory, rule-consequentialism should be preferred over Rossian pluralism. If my argument below is correct, it should show that the first premise of this argument is mistaken. Thus, by their own lights, rule-consequentialists fail to vindicate their view.

Rule-Consequentialism and the Simple Question

Different rule-consequentialists of course formulate their theories in different ways. Generally speaking, according to rule-consequentialists, an act is wrong if it is forbidden by a set of moral principles which, if internalised by the vast majority in the society, would have the best consequences in terms of overall well-being (Brandt 1996, ch. 5; Hooker 1990 and 2000a, p. 32).[3] Otherwise acts are permissible. The set of well-being maximising principles can also require certain acts from us. Not doing these actions is forbidden by the relevant set of principles and therefore wrong.

Before introducing my main argument against rule-consequentialism, I want to set aside a further complication within the rule-consequentialist framework.[4] When we compare which of the sets of moral principles would have the best consequences, we need to consider a set of possible worlds. In each one of these worlds, an alternative set of principles has been internalised. If these worlds are anything like the actual world, the circumstances of these worlds will presumably change from time to time. If we assume that the moral codes of these worlds remain the same throughout, then,

[2] There have also been other attempts to argue that rule-consequentialism has counter-intuitive implications. For instance, it has been argued that rule-consequentialism is too demanding (see Mulgan 2001, ch. 3 for the objection, and Hooker 2004 for a reply). It's not clear whether all these objections are available for the Rossian who also accepts a *prima facie* duty of beneficence.

[3] In general, rule-consequentialists can play with the notion of well-being in order to get the view to fit our intuitions about distributive justice. They can also claim that the well-being of the worse-off counts for more than that of others (Hooker 2000a, pp. 55–9).

[4] I thank the anonymous referee of *Philosophia* for bringing this point up.

in any one world, in some circumstances the adopted code will lead to good lives whereas in other circumstances it will lead to miserable lives. In contrast, if different codes would be internalised in these worlds during different periods, the consequences would be better than the consequences of the stable codes.

However, we have to keep in mind that the people in these worlds must have internalised the relevant moral codes. This makes comparing the worlds in which the moral codes change rather difficult. It's not reasonable to expect the people of these worlds to just suddenly change their moral codes and stable ways of acting. The rules they followed before were after all deeply embedded in their ways of thinking and acting. If this hadn't been the case, the principles would not have worked in the required way.

As a result, in comparing the worlds in which different codes are adopted for different periods, one would also have to take into an account the costs from getting people to change their moral code (Hooker 2000a, pp. 78–80). Given such costs, some circumstances may not be different or enduring enough for it to be worth of it to adopt new codes for them. Yet, in some cases, the benefits of moral reform will have greater expected value than leaving the moral code as it is. Assessing the mechanisms and the costs of the reforms is an interesting and very difficult matter. At some point, the rule-consequentialists would have to address the issue in detail.

I will set these problems aside for now. This is because there is an even more fundamental problem for rule-consequentialism with respect to the change of circumstances. I will therefore assume, for the benefit of rule-consequentialism, that we can compare worlds in which the circumstances remain stable. The circumstances between different subsets of worlds may differ even when they are stable within the worlds. This helps us to compare different sets of moral principles for different circumstances without worrying about the costs of the moral reforms that would eventually have to be taken into account.

The problem I am interested in is not concerned about the change of circumstances in the worlds in which we compare the codes, but rather it concentrates on which code determines the wrongness of actions in new circumstances in the actual, evaluated world. I understand rule-consequentialism here to be a view about the criteria for the wrongness of actions in the actual world. As such, it is neutral about the moral deliberation-procedure we should follow here. For this reason, we do not need to consider how difficult it would be for us to adopt the rule-consequentialist way of thinking or to internalise the ideal code.

To get to the problem I have in mind, we can begin from the subset of worlds in which the circumstances are like they are currently in the actual world. Let us assume that a certain set of moral principles, smp_1, has the best consequences in terms of general well-being in this set of worlds. Imagine then that the circumstances of the actual world change. When we look at the set of worlds in which the circumstances are identical to these new circumstances, it is not smp_1 but rather another, conflicting set of moral principles, smp_2, that has the best consequences in terms of overall well-being.

At this point, we can ask the rule-consequentialists a Simple Question. In the new actual circumstances, are acts wrong when they are forbidden by smp_1, the ideal set for the actual circumstances, or are they wrong when they are forbidden by smp_2, the ideal set for the new circumstances? I will call the version of rule-consequentialism

which gives the first answer 'context-insensitive rule-consequentialism', and the one which gives the second answer 'context-sensitive rule-consequentialism'. I will now argue that both these views lead to normative conclusions which conflict with our considered moral judgments.

Context-Insensitive Rule-Consequentialism

The context-insensitive rule-consequentialists do not change the wrongness-determining moral code for new actual circumstances. We can thus begin by first evaluating the consequences of the alternative moral codes in the worlds in which the circumstances are identical to the current circumstances of our world. Whichever moral code of these worlds has the best consequences then determines which acts are right and wrong in all circumstances. This would match the idea that the moral terms like 'right' and 'wrong' function as so-called rigid designators. That is, some people think that if some types of acts are actually wrong, they must be wrong necessarily in all worlds. This would be the outcome of the context-insensitive rule-consequentialism.

It is obvious that the resulting ideal code requires promise-keeping and forbids promise-breaking.[5] In the actual circumstances, co-operation on contractual basis is possible only if a promise-keeping obligation has been generally internalised. And, without beneficial co-operation much less well-being would be created to the world. The moral code which requires promise-keeping has therefore better over-all consequences in our circumstances than the code which would permit promise-breaking. For this reason, the ideal code for our circumstances presumably requires promise-keeping. If this code determines the rightness and wrongness of actions in all circumstances, then, in the framework of context-insensitive rule-consequentialism, promise-keeping is required in all circumstances.

Imagine then that our circumstances change due to an Evil Demon. In the new circumstances, whenever anyone keeps a promise, some time after this the Demon interferes. As a result of the Demon's devious plots, the over-all well-being of the promisees decreases slightly during a longer period of time. They end up slightly worse off than if the promises had been broken. In these new circumstances, context-insensitive rule-consequentialism still forbids promise-breaking. The case is so constructed that in each case the promisees end up only *slightly* worse off. This means that the rule which requires preventing disasters does not come into play (see Hooker 2000a, pp. 98–9).

However, our considered moral conviction seems to be that, when the actions of the Evil Demon become evident, breaking promises would no longer be morally wrong. One rationale for this may well be that the promise-keeping duty would now diminish the aggregate well-being. Furthermore, no individual would be better off as

[5] This case is from Montague 2000 but the intuitions about it and the conclusions drawn are from Hooker's reply (see Hooker 2000b, sec. 6). Hooker himself uses a similar thought-experiment involving helping others to argue for a similar conclusion. He considers the moral status of altruistic concern. He argues that, in the worlds in which people are very incompetent in helping others, a normal level of such concern would constitute a vice rather than a virtue. This doesn't quite reveal whether he thinks that showing altruistic concern in that world would be morally wrong.

the result of the promise-making practice and its requirements in these circumstances. This seems to imply that the practice and its requirements would collapse because they can no longer serve the practical role they were introduced to serve. Whatever the grounds for our intuitions are, when the moral principles are not changed for the new circumstances, the normative conclusions of rule-consequentialism do not fit our moral convictions about these cases.

Context-Sensitive Rule-Consequentialism

In order to avoid the previous unintuitive results, the rule-consequentialists could change their reply to the Simple Question. This would be to claim that what determines the wrongness of actions in the new circumstance is the code which has the best consequences in the worlds identical to the new circumstances of our world.[6] So, for instance, in the worlds in which there was an evil demon, a code which does not require promise-keeping would have the best consequences. According to context-sensitive rule-consequentialism, this would imply that promise-breaking would not be wrong in our world in the new circumstances. In this section, I want to argue that this view too leads to certain counterintuitive normative conclusions.[7]

I will concentrate on the actions of so-called non-consensual sadism. Sadists are those who can enjoy the suffering of others. In non-consensual sadistic acts, the sadist inflicts pain on others against their wills because she gets pleasure out of their suffering. Most of us think that acts of non-consensual sadism are morally forbidden in *all* possible circumstances (Parfit, unpublished manuscript, ch. 8).

There are several explanations for why we think that such acts are always wrong. A simple Kantian explanation would, for instance, be that the sadist regards her victims as mere means for her selfish ends in an objectionable way. She cares as little about the victims and their strong objections against her actions as she cares about the hammer whilst hammering. Rational human beings, however, deserve always more respect than this. Rule-consequentialists too can say that non-consensual sadism is wrong under the current circumstances because the forbidding rule for

[6] Hooker seems to accept this view (see Hooker 2000a, p. 32, footnote 1; Hooker 2000b, sec. 6; and Hooker 2005, p. 268).

[7] This section is partly based on an observation by McNaughton and Rawling (McNaughton and Rawling 1998, pp. 46–7). They note that, even though the optimific moral principles for the actual world may include intuitive moral constraints, there are possible worlds in which the optimific moral principles do not contain them. So, if these moral constraints were in their absolute form a part of our moral intuitions, then rule-consequentialism would have unintuitive implications. McNaughton and Rawling do not give any examples of such absolute constraints. Hooker's strategy has been to argue that absolute constraints are not part of our moral convictions (Hooker 2000a, pp. 134–6). Philip Montague discusses the earlier circumstances where promise-keeping is generally harmful and therefore keeping promises is forbidden by the optimific principles (Montague 2000). However, our intuition does not seem to be that promise-keeping would be permissible even in those circumstances (see Hooker 2000b, sec. 6). Berys Gaut argues similarly that, based on special relationships, we necessarily have obligations that are independent of their effects on well-being (Gaut 1999, pp. 42–4). These obligations would remain in the worlds in which the relationships might on balance decrease general well-being.

these circumstances has better consequences than the permissive alternative. The victims' pains after all usually outweigh the pleasures of the sadists.[8]

At this point, it could be objected that the rule-consequentialist code cannot even in principle include a rule which forbids non-consensual sadism as an act-type.[9] This is because this type would include a specific motive with which these actions would have to be done in order for them to count as actions of the sort. Acts of sadism must be motivated by the pleasure which the sadist gets from the suffering of her victim.

However, there is nothing in rule-consequentialism itself that prevents us from formulating principles which forbid act-types such that they are in part specified in terms of the motivations with which these actions must be done. Admittedly, the relevant actions must also satisfy some other criteria. The actions done out of happiness do not for example count as an identifiable act-type. But, there are also other substantial criteria which sadistic acts must satisfy besides of being based on a certain motivation. It should be contingent whether the ideal code includes principles which forbid act-types which are in part specified in terms of motivation. The principle which forbids non-consensual sadism presumably does well when we compare the consequences of different principles. It is difficult to imagine a code without this principle which would have good consequences in the current circumstances.

Hooker provides more support for accepting this sort of principles within rule-consequentialism. He argues that we should compare moral codes in terms of the consequences of their general *acceptance* (Hooker 2000a, sec. 3.2). This is because we do not only think about what people do. 'We also care about people's *concerns* (Hooker 2000a, p. 76).' That is, from the point of view of rightness and wrongness, we do not only care about people's actions but also about how they feel and what motivates them to act. If the rule-consequentialist is going to account for the rightness and wrongness of different motivations and feelings, she will have to accept that the compared moral codes can also differ in what kinds of motivations they endorse and discourage. Hooker himself, for instance, compares the consequences of moral codes that prescribe different levels of altruistic concern (Hooker 2000b, sec. 6). Thus, it seems legitimate to assume that a code can forbid act-types which are tied to specific motivations.

The main problem, however, is that there will be new circumstances in which the context-sensitive rule-consequentialism authorises non-consensual sadism with a public moral rule. Consider a world in which an Angel always interferes some time after the sadist has inflicted pain on others against their wills for her own pleasure. The Angel will later on compensate the resulting suffering in some goods conducive to the well-being of the victims and even slightly improve their over-all well-being over the relevant stretch of time. In this scenario, it would be advantageous for the general well-being to have a set of moral principles which publicly authorised non-consensual sadism. After all, the more there was non-consensual sadism the more well-being would be created.

[8] There is an interesting question of whether sadistic pleasures should be taken into account in the cost–benefit analyses of the moral codes at all (see Hooker 2005, p. 276 for a discussion).

[9] I thank the anonymous referee of *Philosophia* for raising this objection.

Yet, in my opinion, it is our considered moral conviction that inflicting pain and suffering on others against their wills for the sake of one's own pleasure would still be wrong even in these new circumstances.[10] Trying to act in this way is wrong, we think, no matter how the consequences of such efforts turn out. This is because these acts continue to show that the sadists consider their victims to be mere tools. The sadists still inflict pain and enjoy the resulting suffering of others. And, more significantly, they still completely ignore the foreseeable objections which their victims would have against their actions. In the Kantian phraseology, their wills are still evil.

Therefore, because context-sensitive rule-consequentialism explicitly supports non-consensual sadism in the new circumstances, the view gives moral permissions which are against our considered moral judgments. These judgments still deem sadism to be wrong in the new circumstances. Of course, the rule-consequentialists can avoid this problem by not changing the wrongness-determining principles for the new circumstances. The idea would then be that the ideal code for our current circumstances determines that non-consensual sadism is wrong also in the new circumstances. This, however, leads to the first horn of the dilemma discussed earlier.

Can Rossian Pluralism Do Better?

The conclusion then is that both context-insensitive and context-sensitive versions of rule-consequentialism lead to normative conclusions which conflict with our moral convictions. We should thus reject rule-consequentialism if we can find another moral theory which coheres better with our moral judgments and if we assess moral theories with the reflective equilibrium test. I want to now argue that there are views that do better in this test than rule-consequentialism.

The arch-rule-consequentialist, Brad Hooker, contends that rule-consequentialism so far fits just as well our considered moral judgments as some forms of Rossian pluralism (Hooker 1996 and 2000a, pp. 106–7). This turns out not to be true if some form of Rossian pluralism can provide the intuitive moral results in both of my thought-experiments. It would also have to be the case that rule-consequentialism cannot use the same Rossian resources to avoid the dilemma described above.[11]

By Rossian pluralism, I mean roughly a view according to which there is an open-ended set of prima facie moral principles which make explicit certain basic and general types of moral reasons. These reasons require, forbid and license actions. I believe that such pluralism can provide the correct answers in both cases above. The pluralist can first accept that there is a distinct prima facie duty not to injure others (Ross 1930, p. 21). We can understand this duty as one which is not limited in its

[10] Notice that the sadist cannot, by definition, aim in the new circumstances at the later well-being of others by doing 'sadistic acts' and hoping that the Angel interferes. This would not count as sadism since the sadist would no longer be motivated solely by the suffering of others but also by their well-being. In such cases, the agent would be aiming at beneficence by pain-inflicting acts. Ironically, the Angel, who is only concerned about pure sadism, would not then interfere and only suffering would be inflicted.

[11] I thank again the anonymous referee of *Philosophia* for pushing this objection.

scope in time. If keeping promises always injures others in the new circumstances after a longer period of time, then we should not keep our promises in order to satisfy this duty. We can say that, in the new circumstances, the duty not to injure others conflicts with and outweighs the duty to keep promises. This is because the latter duty has lost much of its strength in the extraordinary circumstances.

Could a similar method be used to solve to the problem I presented for context-insensitive rule-consequentialism earlier? The plan would be to find such elements from the ideal code for the actual circumstances that, when this code is applied in the new circumstances, it too would no longer require promise-keeping. The first proposal would be to refer to the disaster-prevention rule which is part of the ideal code. This reply would not work. In the new circumstances in which the Demon is present, no disasters are created by promise-keeping. The promisees only end up *slightly* worse off as a result of promise-keeping.

A better option would be to refer to a principle that forbids injuring others. Presumably the ideal code for our circumstances includes such a principle. If we were allowed to injure others, the overall level of well-being in our world would predictably be lower. Thus, the context-insensitive rule-consequentialist could claim that even the ideal code for the old circumstances creates a conflict of duties in the new circumstances in which the Demon is present. The conflicting duties are the principle which requires promise-keeping and the principle which forbids injuring others. The rule-consequentialist could then claim that the latter principle is stronger than the former in the new circumstances, and therefore we are not required to keep our promises when the Demon is present.

The problem with this reply is that the rule-consequentialist cannot be allowed to assign weights for the different principles independently of her theoretical framework. Our moral convictions should not be used within rule-consequentialism to make rule-consequentialism's conclusions match our convictions. If we relied on our convictions at this point, rule-consequentialism would become a mere theoretical 'spare wheel'.[12] Thus, the rule-consequentialist must find a rule-consequentialist rationale for what the weights of the different principles of the ideal set are in different situations.

We can begin from the assumption that, in the worlds we compare the codes, people internalise individual principles with specific degrees of motivational strength (Brandt 1989, p. 95; Hooker 2000a, p. 90). The people of these worlds then solve their conflicts of duties on the basis of the motivational strengths of the principles they have internalised. As a result, we can pick out the ideal code by considering in which of the worlds the motivational strengths of the principles have the best consequences. Finally, when we face actual conflicts of the prima facie principles of the ideal code, the wrongness of our options is determined by which actions would be avoided by the agents who have adopted the ideal motivational strengths of the duties.

Our moral conviction is that, when keeping a promise leaves the promisee only slightly worse off in the current circumstances, we are still required to keep the promise unless the promisee releases us from the duty. This is so unless the promisee would end up *so much* worse off that breaking the promise would count as

[12] See Hooker 2003, p. 58 for a corresponding requirement for contractualism.

preventing a disaster. If we assume that the ideal code for the current circumstances matches this intuition, then the people who have accepted the ideal code for these circumstances must be more motivated to follow the promise-keeping duty than the duty not to injure others slightly.

Context-insensitive rule-consequentialism holds that the ideal code for the actual circumstances determines what is right or wrong in all circumstances. This must also go for how the conflicts between duties are resolved. A given duty is stronger than another one *in all circumstances* if a person who has adopted the ideal code for our circumstances would be more motivated to fulfil that duty in her world than the conflicting one. Therefore, because the people who have internalised the ideal code for our circumstances are more motivated to keep their promises in their world, the promise-keeping duty would outweigh the duty not to injure others slightly even when the Demon is present. This normative consequence, however, still seems unintuitive. The presence of the Demon undermines the promising practice in a way that weakens the strength of the promise-keeping duty. This implies that context-insensitive rule-consequentialism cannot avoid the awkward consequence in the same way as Rossian pluralism.

When it comes to the second horn of the dilemma, the pluralist can also incorporate to her list of duties an independent duty not to aim at our own pleasure by inflicting pain on others against their wills.[13] Given the strength of this duty, the sadist would still act wrongly in the new circumstances. This fits our moral convictions about the new circumstances better than the permission for sadism given by context-sensitive rule-consequentialism.

There is a question about whether context-sensitive rule-consequentialism could use the duty not to injure others discussed above to avoid the awkward permission for sadism. The ideal code for the Angel-including circumstances includes both a prima facie duty not to injure others and a prima facie duty to commit non-consensual sadism. Let us ignore the fact that, even on the prima facie level, the latter duty is highly unintuitive. When we would actually face the circumstances for which this code is ideal, we would still have to make a verdictive judgment about which of these duties is stronger. The context-sensitive rule-consequentialist could try to save the intuitive wrongness of sadism by claiming that, in the new circumstances, the duty not to injure others is stronger than the other duty.

However, again, we cannot let the rule-consequentialist use her moral convictions to determine the strengths of these duties. Rather, the context-sensitive rule-consequentialist has to compare the amount of generated well-being in the Angel-worlds in which individuals have adopted the principles with different motivational strengths. Depending on these motivational strengths, the individuals of these worlds come to solve the conflicts between the duty to not injure others and the duty to do non-consensual sadism in different ways.

The rule-consequentialist's question then is which way of solving these conflicts is most beneficial in terms of general well-being in these worlds. The answer to this question determines which way of acting would be wrong in the corresponding

[13] Ross himself could probably not have accepted such a duty because he did not think that duties can require us to act or not act from specific motives (Ross 1930, pp. 4–6).

actual circumstances according to the context-sensitive view. Presumably, in the worlds in which people have stronger motivation to do non-consensual sadism, the consequences would be better than in the other worlds. After all, every time the sadists overcome their motivation not to injure others, the Angel brings about more well-being.

This means that, according to context-sensitive rule-consequentialism, in the similar circumstances in the actual world, the duty to do non-consensual sadism is stronger than the duty not to harm others. Thus, rule-consequentialists cannot use the duty not to injure others to avoid this horn of the dilemma. The Rossian pluralist, however, can avoid the same unintuitive consequence by adopting a prima facie duty against sadism, and by insisting that the strength of this duty does not depend on its effects on general well-being.

A version of Rossian pluralism can then better cohere with our considered moral intuitions than rule-consequentialism which necessarily fails in one or the other of these cases. The methods with which the Rossian can avoid these problems are not available for the rule-consequentialists, or so I have argued. In the light of the reflective equilibrium test, we should then reject rule-consequentialism. I also believe that there is a more general lesson to be learned from these problems of rule-consequentialism. What the examples show is that we have two kinds of moral convictions—ones that are sensitive to the circumstances and others that are not. If a moral theory can be forced to make a choice about which ones of these intuitions it will be compatible with, the view will be worse off in the reflective equilibrium test than those views that have both context-sensitive and insensitive elements.

References

Brandt, R. (1989). Morality and its critics. *American Philosophical Quarterly, 26*, 89–100.

Brandt, R. (1996). *Facts, values, and morality*. Cambridge: Cambridge University Press.

Daniels, N. (1979). Wide reflective equilibrium and theory of acceptance in ethics. *Journal of Philosophy, 76*, 256–282.

Gaut, B. (1999). Rag-bags, disputes and moral pluralism. *Utilitas, 11*, 37–48.

Hooker, B. (1990). Rule-consequentialism. *Mind, 99*, 67–77.

Hooker, B. (1996). Ross-style pluralism versus rule-consequentialism. *Mind, 105*, 531–552.

Hooker, B. (2000a). *Ideal code, real world*. Oxford: Oxford University Press.

Hooker, B. (2000b). Reflective equilibrium and rule-consequentialism. In B. Hooker, E. Mason, & D. Miller (Eds.), *Morality, rules, and consequences* (pp. 222–238). Edinburgh: Edinburgh University Press.

Hooker, B. (2003). Contractualism, spare wheel, aggregation. In M. Matravers (Ed.), *Scanlon and contractualism* (pp. 53–76). London: Frank Cass.

Hooker, B. (2004). Review of Mulgan's demands of consequentialism. *Philosophy, 78*, 289–296.

Hooker, B. (2005). Reply to Arneson and McIntyre. *Philosophical Issues, 15*, 264–281.

McNaughton, D., & Rawling, P. (1998). On defending deontology. *Utilitas, 11*, 37–54.

Montague, P. (2000). Why rule-consequentialism is not superior to Ross-style pluralism. In B. Hooker, E. Mason, & D. Miller (Eds.), *Morality, rules, and consequences* (pp. 203–211). Edinburgh: Edinburgh University Press.

Mulgan, T. (2001). *The demands of consequentialism*. Oxford: Oxford University Press.

Rawls, J. (1971). *A theory of justice*. Cambridge: Harvard University Press.

Ross, W. D. (1930). *The right and the good*. Oxford: Clarendon.

Williams, B. (1985). *Ethics and the limits of philosophy*. Cambridge: Harvard University Press.