Philosophia (2008) 36:355–366
DOI 10.1007/s11406-007-9115-3

Intentionalism and Change Blindness

Greg Janzen

Received: 6 June 2007 / Revised: 24 November 2007 / Accepted: 6 December 2007 /
Published online: 8 January 2008

Abstract According to reductive intentionalism, the phenomenal character of a conscious experience is constituted by the experience's intentional (or representational) content. In this article I attempt to show that a phenomenon in visual perception called change blindness poses a problem for this doctrine. Specifically, I argue that phenomenal character is not sensitive, as it should be if reductive intentionalism is correct, to fine-grained variations in content. The standard anti-intentionalist strategy is to adduce putative cases in which phenomenal character varies despite sameness of content. This paper explores an alternative antiintentionalist tack, arguing, by way of a specific example involving change blindness, that content can vary despite sameness of phenomenal character.

Keywords Intentionalism · Change blindness · Consciousness · Intentional content

Introduction

Intentionalism is the view that the phenomenal character of a conscious experience supervenes on the experience's intentional (or representational) content, such that, necessarily, "there can be no difference in phenomenal character without a difference in content" (Byrne 2001, p. 204).[1] Thus, if two experiences differ in phenomenal character, then they differ in intentional content (hereafter simply "content"). Reductive intentionalism, as I will call it, makes a stronger claim, namely, that an experience's

[1]By "phenomenal character" I take intentionalists to mean, roughly, "appertaining to what a conscious state is *like* from the first person perspective." Here is Byrne, a prominent intentionalist, on the matter: "The notion of the phenomenal character of an experience is hard to explain, but easy to understand. (At any rate everyone seems to understand it.) We can start with the stock phrase: 'what it's like' for the subject to undergo the experience. We can give everyday examples of similarity and difference in phenomenal character: the experience of seeing purple is more like, in respect of phenomenal character, the experience of seeing blue than it is like the experience of smelling vanilla" (2001, p. 200; also see Tye 1995, p. 3).

G. Janzen (✉)
Department of Philosophy, University of Calgary, Calgary T2N1N4, Canada
e-mail: gajanzen@ucalgary.ca

phenomenal character is *constituted by* its content. On this stronger version, content fully exhausts the nature of phenomenal character. This view is widely held nowadays,[2] especially by those seeking a naturalistically acceptable account of consciousness. The idea is that if an experience's phenomenal character is constituted by its content, then, since it may be possible to explain content in naturalistic terms, there is at least the prospect of a wholly naturalistic theory of consciousness.

Reductive intentionalism is deeply problematic, however, and so while the search for a naturalistic theory of consciousness is laudable, those seeking such a theory may have to abandon the doctrine in favour of some other piece of naturalizing machinery. Or so I shall argue. My aim is to dispute a largely unchallenged assumption on which the view is based, namely, that, in addition to *course-grained* or *inter-modal* differences in phenomenal character – e.g., the difference between what it is like to *see* an apple and what it is like to *feel* an apple – there are *fine-grained* or *intra-modal* differences in phenomenal character – e.g., the difference between what it is like to see a *red apple* and what it is like to see a *green apple* (or the difference between what it is like to see an apple at *t* and what it is like to see the same apple at *t*+, where, between *t* and *t*+, the apple has undergone a perceptible surface change).[3] Phenomenal character, I want to suggest, is not sensitive in this way to fine-grained variations in content. The usual anti-intentionalist strategy is to adduce putative cases in which phenomenal character varies despite sameness of content.[4] This paper explores an alternative anti-intentionalist tack, arguing that content can vary despite sameness of phenomenal character. My strategy is to present three cases involving a phenomenon in visual perception called *change blindness*. I argue that the reductive intentionalist can plausibly rebut the first two cases, but not the third. It follows from the success of this case that content can vary independently of phenomenal character. But if content can vary independently of phenomenal character, then it is not the case that phenomenal character is constituted by content, and so reductive intentionalism is false.

Reductive Intentionalism

My interest in this paper is with reductive intentionalism, a theory that, as Tye puts it, "aims to tell us what phenomenal character *is*" (2000, p. 45). For ease of exposition, I will hereafter refer to the view simply as "intentionalism." There are different versions of the theory, depending on intentionalism's scope. *Unrestricted* intentionalism is the thesis that intentionalism applies to *all* conscious experiences, whereas *restricted* intentionalism leaves open the possibility that there are certain experiences whose phenomenal character is not constituted by their content. Dretske (1995) has defended a kind of restricted intentionalism, arguing that although the phenomenal character of perceptual experiences is constituted by their content, this may not be the case for the phenomenal character of moods and emotions. Tye, on the other hand, has defended a kind of unrestricted intentionalism. On his view, the phenomenal character of *all*

[2] See, e.g., Crane 2000; Dretske 1995, 2003; Lycan 1996; Rey 1998; Tye 1995, 2000.

[3] Many extend this view to purely cognitive states, arguing that what it is like to think that *p* is different from what it is like to think that *q* (e.g., Woodruff Smith 1989, Chalmers 1996, Pitt 2004, Zahavi 2005).

[4] Peacocke (e.g., 1983) and Block (e.g., 1996), among others, have advanced cases of this sort.

conscious experiences is constituted by their content (see Tye 1995, chs. 4 and 5). Now, some have argued that restricted intentionalism is problematic. According to Byrne (2001, pp. 205–6), for example, it is unclear how one might simultaneously show that intentionalism holds in the case of, say, visual experiences but fails in the case of, say, itches (assuming that itches, like visual experiences, have intentional content); and Kind (2007) argues that intentionalism has to be unrestricted to serve as an adequate theory of qualia. This issue, however, need not detain us unduly, since my aim is to show that intentionalism fails where it purports to be on its firmest ground, i.e., in the case of perceptual experience. And if the doctrine fails in the case of perceptual experience, then it probably fails in the case of pains, itches, tickles, and the like. Of course, it is *possible* that, despite the failure, as I intend to show, of intentionalism regarding perceptual experience, an intentionalist could make a case for intentionalism regarding, say, pains. But it is altogether unclear how that could be done.

Change Blindness: Three Cases

Both proponents and opponents of intentionalism typically assume that the phenomenal character of perceptual experience is sensitive to the distinguishable surface qualities of perceived objects, so that, necessarily, if a perceived surface quality changes (i.e., if there is a change in content), then the phenomenal character of one's perceptual experience of that surface quality changes. Tye is explicit on this:

> As you view the scene before your eyes and how things look to you, necessarily, if any of the qualities of which you are *directly* aware change, then the phenomenal character of your experience changes. Consider, for example, the facing surface of a ripe tomato in a bowl before you. In attending to the color of the tomato, you are directly aware of a certain quality, *Q*, as covering that surface. You experience each just noticeable part of the surface as having *Q*... [W]hether *Q* is itself a color or some other quality, awareness of which mediates your awareness of color, is left open here. But change *Q* – for example, by changing the color of the tomato or by donning color-inverting lenses – and what it is like for you in viewing the tomato necessarily changes. (2000, p. 48)

The assumption that variations in content are sufficient for variations in phenomenal character has met with little resistance. To be sure, Block (e.g., 1990, pp. 61ff.) has described a case that is the "converse" of the usual inverted spectrum cases, i.e., a case of inverted intentional content combined with identical phenomenal content. But his case relies on certain controversial assumptions – e.g., that inverted spectra are possible – that I am keen to avoid. In contrast to inverted spectrum scenarios, my cases appeal to a bizarre yet relatively oft-occurring visual phenomenon that Rensink et al. (1997) have dubbed *change blindness*, which occurs when a normally sighted subject has difficulty detecting (sometimes significant) changes within her visual field when the changes are viewed successively.[5]

[5] Change blindness (or *difference blindness* as I will call it) is to be distinguished from a related visual phenomenon called *inattentional blindness*, which is the inability to perceive (sometimes obvious) features in a visual scene when the subject is not attending to them.

Following Dretske (2004), I will speak of *difference blindness* rather than change blindness, because in situations in which change blindness occurs, subjects are exposed to differences but do not experience the change that causes them. In fact, in experiments used to demonstrate change blindness, the change is purposely concealed. "Difference blindness" is therefore more appropriate, since subjects are blind to differences, not changes (which are not visible at all). I will state the cases and then argue that the third constitutes a counterexample to intentionalism.

(C1) Suppose S, a competent adult with no cognitive deficiencies, looks at a photograph of twenty red roses at t and then looks at the same photograph at $t+$, failing to notice at $t+$ that, between t and $t+$, one of the roses was airbrushed out of the photograph. Since S saw a photograph of twenty roses at t and a photograph of nineteen roses at $t+$, her visual experience at t differs from her visual experience at $t+$. Nevertheless, S is not *aware* that her experiences differ; and even after giving the matter some additional thought, she denies that she saw any difference in the photograph.

(C2) Suppose there's a grapefruit on S's kitchen counter that she sees every time she enters her kitchen. She enters her kitchen about a dozen times a day, and so she sees the grapefruit about a dozen times a day. Further suppose that as the grapefruit ripens, its colour changes from yellow to yellow* (a different shade of pale yellow). The change would be noticeable if it was to occur quickly, but since the change occurs gradually, S never notices it. She enters her kitchen one morning and eats the grapefruit, never knowing that it underwent a colour change.

(C3) Dretske (2004), via Grimes (1996, p. 92), describes an example of difference blindness documented by McConkie and Zola (1979). Experimental subjects were asked to read the following text in which the letters are alternated between upper and lower case:

(1) ThE sPaCe ShUtTlE tHuNdErEd InTo ThE sKy On A cOlUmN oF sMoKe

The subjects' eye movements were monitored and during each saccade (a ballistic eye movement that occurs when the subject fixes her eyes on one point after another in the visual field) the case of every letter was changed. That is, (1) was replaced by (2):

(2) tHe SpAcE sHuTtLe ThUnDeReD iNtO tHe SkY oN a CoLuMn Of SmOkE

Then, during the next saccade, (2) was replaced by (1). However, as Dretske puts it:

> The alternation continued so that the appearance of every letter was dramatically different on each new fixation. Since the change took place during a saccade when the eye transmits no information, subjects did not see the change (event) that occurred several times every second. That was no surprise. What was surprising is that subjects failed to notice a difference in what they saw on each fixation – the difference between (1) and (2). There was no awareness that anything unusual was going on. Although every letter they saw changed its appearance several times each second, they never noticed it. (2004, p. 6)

Intuitions will no doubt diverge on this, but it is plausible to suppose that, in each of these cases, the phenomenal character of the subject's visual experience at *t* does not differ from the phenomenal character of her visual experience at *t*+, whereas the content of her experience does. In C_1, for example, since S saw a photograph of twenty roses at *t* and a photograph of nineteen roses at *t*+, her visual experience at *t* differs in content from her visual experience at *t*+: at *t* her visual experience was of a photograph of twenty roses, whereas at *t*+ it was not. Yet phenomenally her experiences do not differ. The same can be said, *mutatis mutandis*, for C_2 and C_3. It would seem to follow that content can vary independently of phenomenal character.

My strategy for defending C_3 as a counterexample to intentionalism will be to consider two possible ways for an intentionalist respond to C_1–C_3 and show that, when it comes to C_3, neither is convincing. Pursuing this strategy is appropriate because I think it is *prima facie* plausible that in all three cases the subject's successive visual experiences represent the world as being a different way even though the phenomenal character of her experiences remains constant. It is best to proceed, therefore, by considering attempts to show wherein the cases fail.

Intentionalist Responses Rebutted

There are two possible ways for an intentionalist to respond to C_1–C_3: (1) argue that, despite appearances, there really is no difference in content, or (2) argue that there is a difference in phenomenal character. I will examine these responses in turn.

To defend intentionalism on the grounds that there is no difference in content, an intentionalist might appeal to the distinction between *personal* and *sub-personal* levels of content. There are a number of complicated questions surrounding the notion of content – e.g., whether it is best treated along Russellian, Fregean, or possible worlds lines;[6] whether it ought to be individuated widely in a way that is sensitive to social and environmental factors; and so on – but all the disputants in the literature agree that content "specifies the way the world appears or seems *to the subject*" (Byrne 2001, p. 201). The basic idea is that the relevant content is personal-level (or conscious, or experiential) content. Call this kind of content "PL-content." Consider the famous Müller-Lyer illusion, where a subject is presented with two lines of equal length, one with arrows directed outward and the other with arrows directed inward. Even though the lines are of equal length, the content of the subject's visual experience is of two lines of unequal length, since *to the subject* the line with the arrows directed outward appears longer than the line with the arrows directed inward. Intentionalists countenance sub-personal content (e.g., implicit representations), but since phenomenal character is, incontestably, a personal-level phenomenon – it makes no sense to say, for example, that the phenomenal character of S's perceptual experience of x occurs at a sub-personal level – it is, if intentionalism is correct, constituted by PL-content. In response to C_1–C_3, then, an intentionalist might claim that, *at the personal level*, the content of S's visual experience at *t* does not differ in content from her visual

[6] For discussion of these treatments of content see, respectively, Peacocke (1992), Woodruff Smith (1989), and Stalnaker (1999).

experience at *t*+. That is, even though the sub-personal content of her experience at *t* may differ from the sub-personal content of her experience at *t*+ , it does not follow that the PL-content differs: the way the world appears visually *to the subject* is the same at both *t* and *t*+. But since phenomenal character is constituted by PL-content, C_1–C_3 do not jeopardize intentionalism.

This response might be apt when it comes to C_1: it seems accurate to affirm that, at the personal level, the content of S's experience is indeterminate, that what is represented to the subject is a photograph of a *bunch* of roses. Indeed, when asked how many roses she saw (at both *t* and *t*+) S would say "Many" or "Several," and this is the natural, indeed only, response to be given here (assuming, of course, that S did not count the number roses in the photograph). So the sub-personal content of S's visual experience at *t* is (in part) a photograph of twenty roses, whereas the sub-personal content of her visual experience at *t*+ is (in part) a photograph of nineteen roses. However, the PL-content of S's visual experience at both *t* and *t*+ is a photograph of a bunch of roses.

This response seems entirely unapt, however, when it comes to the other two cases. Consider C_2. It strikes me as virtually platitudinous that the way the grapefruit appears *to the subject* at *t* differs from the way the grapefruit appears to the subject at *t*+ (i.e., on the day she eats it): at *t*, the grapefruit appears to the subject as a *yellow* grapefruit, whereas at *t*+ the grapefruit appears to the subject as a *yellow** grapefruit. In C_3, similarly, the PL-content of the subject's visual experience at *t* is this:

(1) ThE sPaCe ShUtTlE tHuNdErEd InTo ThE sKy On A cOlUmN oF sMoKe,

whereas the PL-content of the subject's visual experience at *t*+ is apparently quite different, namely, this:

(2) tHe SpAcE sHuTtLe ThUnDeReD iNtO tHe SkY oN a CoLuMn Of SmOkE.

Indeed, to suppose otherwise is to suppose, *prima facie* implausibly, that the PL-content of S's visual experience when she sees (2) is identical to the PL-content of S's visual experience when she sees (1).

This is too quick, however, for there is a model for understanding difference blindness – what Dretske calls the *object* model – that may have the resources required to make a case for the claim that in both C_2 and C_3 there is in fact no difference in PL-content between *t* and *t*+.[7] According to the object model, subjects who claim to see no difference are probably wrong: they do see it; they just don't notice it. If the object model is correct, then the subject in C_2 and C_3 sees the difference; she just doesn't notice it. Dretske compares differences on the object model to spies: just as one can see a spy (e.g., one's neighbour) and not notice that one has seen a spy, so one can see a difference and not notice that one has seen a difference. The object model is to be distinguished from what Dretske calls the *fact*

[7] There are some difficult questions regarding difference blindness. For example, what is the role of implicit representation in difference blindness? Is attention required to perceive differences? Is it important for subjects to *actively search* for differences in reducing difference blindness? For recent discussions of these questions see, respectively, Silverman and Mack (2006), Beck et al. (2007), and Rensink et al. (1997). My argument in this paper does not demand that I take up a position on any of these controversies.

model, according to which difference blindness is an epistemic or cognitive failure, i.e., a blindness to facts – e.g., the fact that there is a difference between the yellow grapefruit and the yellow* grapefruit. On this model, one cannot see a difference without noticing it, since to see a difference is to see a fact, and "facts (e.g., that x has changed color), unlike concrete objects (x) and properties (colors), cannot be seen without being noticed" (ibid.). Another important (at least for our purposes) feature of the fact model is that, unlike on the object model, subjects can be "blind to differences and still see all the things (objects and their properties) that make up the difference" (Dretske 2004, p. 11). Thus, if the fact model is correct, the subject in C_2 and C_3 does not see the difference, even though she sees all the objects and properties that make up the difference.

As noted above, an intentionalist might appeal to the object model to explain how it is that in C_2 and C_3 there is no difference in PL-content between *t* and *t*+. Specifically, it might be argued that there is no difference in PL-content because even though the subject sees the difference, she doesn't notice it, and hence her visual awareness of it takes place at the sub-personal level. In C_3, for example, the PL-content of the subject's experience when she sees (1) is a series of Gestalt groupings of letters, and the PL-content of her experience when she sees (2) is a series of Gestalt groupings that is identical to the series of Gestalt groupings that she sees when she sees (1). Thus, to the subject, the world appears the same at both *t* and *t*+, and hence there is no difference in PL-content.

The problem with this response, however, is that the object model for understanding difference blindness is eminently disputable. A difference isn't any kind of object or property. It makes sense to say that x has the property of being different from y, but the property being ascribed to x in this case is a *relational* property, the property of being related to y differentially. But the proponent of the object model is committed to affirming that, in failing to notice a difference, one fails to notice some kind of *intrinsic* property (or some kind of object). But x's being different from y is no more an intrinsic property of x than x's being taller than y is an intrinsic property of x.

To see a difference between x and y, then, is not to see any kind of object or property, but rather a fact, i.e., the fact that x and y are different. If, for example, S fails to see the difference between the yellow grapefruit and the yellow* grapefruit, she fails to see a fact, the fact that the grapefruit changed from yellow to yellow*. But the fact model of understanding difference blindness will not help the intentionalist. If the fact model is correct, then the subject in both cases does not see any differences, but still sees all the properties and objects that make up the differences. Thus, in C_3, the subject does not see the difference between (1) and (2), but sees the objects and properties that make up the difference. But if she sees the objects and properties that make up the difference, then at *t* she sees – consciously sees – this:

(1) ThE sPaCe ShUtTlE tHuNdErEd InTo ThE sKy On A cOlUmN oF sMoKe

and at *t*+ she sees this:

(2) tHe SpAcE sHuTtLe ThUnDeReD iNtO tHe SkY oN a CoLuMn Of SmOkE.

But if she sees (1) at *t* and (2) at *t*+, then her experience at *t* differs in PL-content from her experience at *t*+. The reason is that (1) makes up part of the PL-content of her experience at *t*, whereas (2) makes up part of the PL-content of her experience at

t+, and (1) and (2) are represented differently.[8] Exactly parallel considerations apply to C_2. If the subject sees the objects and properties that make up the difference between the yellow grapefruit and the yellow* grapefruit, then at *t* she sees the yellow grapefruit and at *t*+ she sees the yellow* grapefruit. But then the PL-content of her visual experience at *t* must differ from the PL-content of her visual experience at *t*+, since at *t* the PL-content of her experience is constituted (in part) by a yellow grapefruit, whereas at *t*+ it is not.

It seems, then, that the proponent of intentionalism will have to defend the doctrine against C_2 and C_3 by arguing that S's visual experience at *t* differs phenomenally from her visual experience at *t*+. To do this, an intentionalist might argue that, in both C_2 and C_3, there is a difference in phenomenal character between *t* and *t*+; only, the difference is so small that the subject doesn't notice it. This response would, I expect, find an ally in Byrne, who claims that "[b]ecause the subject...may have limited powers of discrimination (like us), it is a mistake to hold that she will always *know* that there is a change in phenomenal character: if the change is sufficiently small, she won't" (2001, p. 209n19). The assumption motivating this remark is that if the perceptible qualities of an object have changed *negligibly*, then a subject may not notice – and hence not *know* (since, at least in perception arguably, noticing that *p* is a necessary condition for knowing that *p*) – that phenomenally her experience prior to the change was different from her to experience after the change.[9] It might be claimed, on the basis of this assumption, that the subjects in C_2 and C_3 are simply unaware that their successive visual experiences differ phenomenally.

This sort of response, however, seems unjustifiably stipulative, especially in light of the phenomenology. Consider C_3. There, unlike in C_2, the difference in what is represented is quite conspicuous; indeed, the difference is a *big* difference, the kind that surprises and nonplusses subjects when they realize the obviousness of it. And so it strains credibility to suggest that there is a difference in phenomenal character between *t* and *t*+ that has gone unnoticed. Put another way, according to intentionalism, phenomenal character is constituted by content, and so any conspicuous difference in content should, if intentionalism is correct, result in a noticeable difference in phenomenal character. But in C_3 it doesn't.

[8] One might object that the individual letters that make up (1) and (2) are not part of the PL-content of the subject's visual experience, since the subject does not attend to them. (Consider Tye's claim that "necessarily, if any of the qualities of which you are *directly* aware change, then the phenomenal character of your experience changes" (2000, p. 48)). This objection misfires, however, because it commits one to affirming that only those objects and properties of an experience that are attended to can form the PL-content of one's experience, which is implausible. Surely the blue sky, of which I am now peripherally aware as I look outside my window at the mountain in the distance, constitutes part of the PL-content of my visual experience of the mountain. Similarly, the individual letters in (1) and (2), even if the subject does not attend to them, form part of the PL-content of the subject's visual experience of (1) and (2).

[9] Byrne makes this assumption because although he thinks changes in phenomenal character are self-intimating, he finds compelling an argument from Williamson (2000, ch. 4) against luminosity, which Williamson defines as the thesis that for a condition C and for any case α, one is in a position to know that C obtains if and only if, in α, C obtains. The idea is that if C is luminous, then one will know, in any case in which it obtains, that it has in fact obtained. In his anti-luminosity argument, Williamson uses the example of feeling cold, arguing that, contrary to accepted belief, feeling cold is a condition that is not luminous at all.

That the difference in content is so conspicuous in C_3 should also incline us to be less than sanguine about the possibility that the alleged difference in phenomenal character is combined with near-instantaneous memory loss. The possibility of such "phenomenal amnesia" cannot be ruled out a priori, but if there is a difference in phenomenal character between *t* and *t*+, then, given that there is a relatively conspicuous difference in content, there should be an equally (or at least somewhat) conspicuous difference in phenomenal character, one that isn't immediately forgotten. Although a plausible response, then, to C_2 is to say that there is a difference in phenomenal character between *t* and *t*+ that has gone unnoticed, such a response fails against C_3.

Still, an intentionalist will insist that, at most, we have reached a dialectical impasse: the intentionalist says that there is a difference in phenomenal character in C_3 that has gone unnoticed, while the ant-intentionalist says there is no difference at all. But since the anti-intentionalist's case is not decisive, it falls short as a counterexample to intentionalism.

There is a consideration, however, that tips the scales in favour of the view that there is no difference in phenomenal character, namely, the subject's sincere phenomenological report of her successive visual experiences. Block expresses the intuition here as follows: "The phenomenal character of your…experience stays the same. That's what you say, and why shouldn't we believe you?" (1996, p. 42). Indeed, for how long, and on what basis, could the intentionalist insist on the charge that the subject's successive visual experiences in C_3 differ phenomenally? It seems a fallback position that unless strong reason can be found for repudiating the phenomenological testimony of reliable subjects – and I know of no such reason – the claim that there is no difference in phenomenal character between *t* and *t*+ goes through cleanly. This is just to call attention to an important feature regarding the appearances in question that should be borne in mind. That there is no difference in phenomenal character between *t* and *t*+ has phenomenology firmly on its side: the subjects in C_3 insist that there is no difference. And, unlike in C_2, the results of such phenomenological reflection cannot be explained away in terms of a difference in phenomenal character that has gone unnoticed.

An objection that might be raised here is that a "suitably idealized" subject (cf. Byrne 2001), one who, given her particularly acute powers of phenomenal discrimination, would notice *any* difference in phenomenal character, and hence would notice the difference in phenomenal character between *t* and *t*+. However, this objection presupposes that there is a difference in phenomenal character between *t* and *t*+, which his precisely what is being denied. The claim is not that there is a difference in phenomenal character that the subject does not notice; the claim is, rather, that there is no difference in phenomenal character to *be* noticed. Of course, if the subject in C_3 notices the *change* in content from *t* to *t*+, then the subject would doubtless notice the *difference* in content between (1) and (2); and a difference in the phenomenal character of the subject's successive visual experiences of (1) and (2) would thereby result. But this consideration is quite beside the point, since, *in the case described*, the subject does not notice the change in content. Thus, to stipulate that the subject notices the change is an evasive manoeuvre that neglects to deal with the case on its own terms.

The foregoing privileging of phenomenology is not without principled support. Underlying it is the independently motivated assumption that there is no appearance/

reality distinction when it comes to conscious experience.[10] A fallacy that has become distressingly prevalent in recent attempts to understand consciousness involves supposing that the way a conscious experience is in itself might transcend the way it is for the experiencing subject, i.e., the way phenomenology reveals it to be.[11] For example, in response to Edelman's claim that "[o]ne of the most striking features of consciousness is its continuity" (1989, p. 119), Dennett says:

> This is utterly wrong. One of the most striking features about consciousness is its discontinuity – as revealed in the blind spot, and saccadic gaps, to take the simplest examples. The *dis*continuity of consciousness is striking because of the *apparent* continuity of consciousness. (1991, p. 356)

This claim has dramatic implications: if Dennett is right, then perceptual consciousness, as it really is, is irregular and sporadic. The continuity we experience in naïve perception, one is given to understand, is merely "apparent," an illusion. It follows that conscious perceptual experience is one way in itself and one way for the experiencing subject; and the way it *really* is, i.e., the way it is in itself, is often quite distinct from how it is for the experiencing subject.

But this is profoundly misconceived. To suggest that consciousness, perceptual or otherwise, might transcend the way it is for the experiencing subject betokens a failure to heed the distinction, previously considered, between personal and the sub-personal levels of content. At the sub-personal level, we have various neurological (cognitive, etc.) processes that are, presumably, necessary for experience. At the personal level, by contrast, we have the "field of experience": this is the level of content that phenomenology describes, the level of content that we experience from the first-person perspective.[12] And no sub-personal consideration can displace or supersede this level of content. Thus, it makes no sense to say that perceptual consciousness, despite its apparent continuity, is *really* marked by discontinuity, just as it makes no sense to say that pain, despite its apparent painfulness, is not *really* painful. There is no such appearance/reality distinction to be made with respect to conscious experience. As Searle has so perspicuously put it: "where the existence of conscious states is concerned, you can't make the distinction between appearance and reality, *because the existence of the appearance is the reality in question*" (1997,

[10] I cannot provide a responsible defence of this assumption within the ambit of this paper (below I provide a mere statement of it), so I will simply defer to others who have. Various classical phenomenologists, including Brentano (1874), Sartre (1956), and Merleau-Ponty (1962), have defended it, but a number of contemporary philosophers would also ally themselves with it. See, e.g., Nagel (1974), Woodruff Smith (e.g., 1989), Dwyer (1990), Siewert (1998), Stubenberg (1998), Thomasson (2000), Zahavi (e.g., 1999, 2005), and Carman (2005).

[11] For interesting and stimulating recent discussions of this fallacy see Dwyer (1990) and Rowlands (2001). Both Dwyer and Rowlands characterize the fallacy as a category mistake. According to Dwyer, subjectivity encompasses "categorially autonomous phenomena" (1990, p. 33), and so it is a fallacy to suppose that our actual experience "may be some way in itself transcending the way it is for a subject" (1990, p. 32); and Rowlands claims that since personal and sub-personal levels of content cannot be reckoned to be in the same category, it is a fallacy to suppose that there is a "distinction between (i) the way an experience seems to its subject and (ii) the way an experience really is" (2001, p. 189).

[12] The wording here is perhaps somewhat infelicitous. Strictly speaking, we do not experience a level of content; rather, we experience objects, events, properties, states of affairs, etc. The level of content described in phenomenology just is our experience of these worldly items.

pp. 121–2).[13] This is not to deny that the neurological (cognitive, etc.) structures that are causally related to conscious experience (e.g., the neurological structures that are indispensable for seeing blue) merit investigation; and this is not to deny that we should endeavour to provide a phenomenological framework in which the phenomenological properties recognized are made continuous with the properties admitted by neuroscience, cognitive science, etc. However, it is to deny that sub-personal considerations somehow take precedence in the study of consciousness, i.e., that sub-personal considerations can give us insight into the "real" nature of consciousness, whereas the experiential considerations phenomenology describes cannot. Phenomenology, in other words, antecedes the empirical sciences of the mind: the data they disciplines provide must be made to cohere with the phenomenological data, not vice versa. One upshot here is that if there is a standoff between two accounts of the phenomenal character of an experience, but one involves denying the subject's sincere phenomenological report of the experience while the other does not, surely one should opt for the latter.

It bears emphasizing that it does not follow from the fact that there is no appearance/reality distinction to be made with respect to conscious experience that one cannot be mistaken about the object or, more generally, the content of one's experience. It is obviously true that sometimes we err in our perceptual judgements, thinking we saw (heard, felt, etc.) something when we saw (heard, felt, etc.) something else. But all that follows from this is that sometimes we lack the information necessary to make a sound judgement about the object or content of our experience, or that sometimes our perceptual capacities at the experiential level outstrip our conceptual/cognitive resources. The thesis under consideration is that there is no appearance/reality distinction to be made with respect to conscious experience – i.e., that one cannot be mistaken about how a conscious experience *seems* or *appears* to one – a thesis that is not impugned by the obvious fallibility of our perceptual judgements. I can be visually experiencing a stick in water and mistakenly think it is bent, but I can hardly be mistaken about it seeming to me as though I am visually experiencing a bent stick in water.

Conclusion

I conclude that C_3 is immune to available intentionalist responses. As anyone even vaguely familiar with the pertinent literature well knows, it is exceedingly difficult to make a compelling case against intentionalism merely by dint of counterexample. In addition, it might be thought that it is futile to attempt to draw any substantive philosophical conclusions from the cases on which I have expounded, since the cases express intuitions about quintessentially empirical questions, questions to which scientific answers are now only forming. However, C_3 is, I think, at least *suggestive*: it poses a genuine problem for intentionalism, and therefore puts the ball back in the intentionalist's court.

[13] See also Honderich (2004, pp. 132–3), McGinn (1995, p. 247), Nagel (1974, p. 448), and Stubenberg (1998, p. 38).

References

Beck, M., Levin, D., & Angelone, B. (2007). Change blindness blindness: Beliefs about the roles of intention and scene complexity in change detection. *Consciousness and Cognition*, *16*, 31–51.

Block, N. (1990). Inverted earth. *Philosophical Perspectives*, *4*, 53–79.

Block, N. (1996). Mental paint and mental latex. *Philosophical Issues*, *7*, 19–49.

Brentano, F. (1874). *Psychology from an empirical standpoint*. L. L. McAlister (Ed.). A. C. Rancurello, D. B. Terrell, and L. L. McAlister (Trans.). (London: Routledge, 1973).

Byrne, A. (2001). Intentionalism defended. *Philosophical Review*, *110*, 199–240.

Carman, T. (2005). On the inescapability of phenomenology. In D. Woodruff Smith, & A. Thomasson (Eds.) *Phenomenology and the philosophy of mind* (pp. 67–89). Oxford: Oxford University Press.

Chalmers, D. (1996). *The conscious mind: In search of a fundamental theory*. Oxford: Oxford University Press.

Crane, T. (2000). Introspection, intentionality, and the transparency of experience. *Philosophical Topics*, *28*, 49–67.

Dennett, D. (1991). *Consciousness explained*. Boston: Little, Brown and Company.

Dretske, F. (1995). *Naturalizing the mind*. Cambridge, MA: MIT Press.

Dretske, F. (2003). Experience as representation. *Philosophical Issues*, *13*, 67–82.

Dretske, F. (2004). Change blindness. *Philosophical Studies*, *120*, 1–18.

Dwyer, P. (1990). *Sense and subjectivity: A study of Wittgenstein and Merleau-Ponty*. Leiden: E. J. Brill.

Edelman, G. (1989). *The remembered present: A biological theory of consciousness*. New York: Basic Books.

Grimes, J. (1996). On the failure to detect changes in scenes across saccades. In K. Akins (Ed.) *Perception* (pp. 89–110). Oxford: Oxford University Press.

Honderich, T. (2004). *On consciousness*. Edinburgh: Edinburgh University Press.

Kind A. (2007). Restrictions on representationalism. *Philosophical Studies*, *134*, 405–427.

Lycan, W. G. (1996). *Consciousness and experience*. Cambridge, MA: MIT Press.

McConkie, G. W., & Zola, D. (1979). Is visual information integrated across successive fixations in reading? *Perception and Psychophysics*, *25*, 221–224.

McGinn, C. (1995). Consciousness evaded: Comments on Dennett. *Philosophical Perspectives*, *9*, 241–249.

Merleau-Ponty, M. (1962). *Phenomenology of perception*. Colin Smith (Trans.). London: Routledge.

Nagel, T. (1974). What is it like to be a bat? *Philosophical Review*, *83*, 435–450.

Peacocke, C. (1983). *Sense and content*. Oxford: Oxford University Press.

Peacocke, C. (1992). *A study of concepts*. Cambridge: MIT Press.

Pitt, D. (2004). The phenomenology of cognition, or what is it like to think that p? *Philosophy and Phenomenological Research*, *69*, 1–36.

Rensink, R., O'Regan, J., & Clark, J. (1997). To see or not to see: The need for attention to perceive changes in scenes. *Psychological Science*, *8*, 368–373.

Rey, G. (1998). A narrow representationalist account of qualitative experience. *Philosophical Perspectives*, *12*, 435–457.

Rowlands, M. (2001). *The nature of consciousness*. Cambridge: Cambridge University Press.

Sartre, J. P. (1956). *Being and nothingness*. H. Barnes (Trans.). New York: Philosophical Library.

Searle, J. (1997). *The mystery of consciousness*. New York: New York Review of Books.

Siewert, C. (1998). *The significance of consciousness*. Princeton: Princeton University Press.

Silverman, M., & Mack, A. (2006). Change blindness and priming: When it does and does not occur. *Consciousness and cognition*, *15*, 409–22.

Stalnaker, R. (1999). *Context and content*. Oxford: Oxford University Press.

Stubenberg, L. (1998). *Consciousness and qualia*. Amsterdam/Philadelphia: John Benjamins Publishing Company.

Thomasson, A. (2000). After Brentano: A one-level theory of consciousness. *European Journal of Philosophy*, *8*, 190–209.

Tye, M. (1995). *Ten problems of consciousness: A representational theory of the phenomenal mind*. Cambridge, MA: MIT Press.

Tye, M. (2000). *Consciousness, color, and content*. Cambridge, MA: MIT Press.

Williamson, T. (2000). *Knowledge and its limits*. Oxford: Oxford University Press.

Woodruff Smith, D. (1989). *The circle of acquaintance: Perception, consciousness, and empathy*. Dordrecht: Kluwer.

Zahavi, D. (1999). *Self-awareness and alterity: A phenomenological investigation*. Evanston: Northwestern University Press.

Zahavi, D. (2005). *Subjectivity and selfhood: Investigating the first-person perspective*. Cambridge, MA: MIT Press.