

‘Is’–‘Ought’ Derivations and Ethical Taxonomies

Scott Hill

Received: 7 December 2007 / Revised: 28 March 2008 / Accepted: 1 April 2008 /
Published online: 17 April 2008
© Springer Science + Business Media B.V. 2008

Abstract Hume seems to claim that there does not exist a valid argument that has all non-ethical sentences as premises and an ethical sentence as its conclusion. Starting with Prior, a number of counterexamples to this claim have been proposed. Unfortunately, all of these proposals are controversial. Even the most plausible have a premise that seems like it might be an ethical sentence or a conclusion that seems like it might be non-ethical. Since it is difficult to tell whether any of these counterexamples are genuine, we need a taxonomy that sorts out ethical sentences from non-ethical ones. We need to know the difference between an ‘Is’ and an ‘Ought’. In the first part of the paper, I establish the need for a taxonomy. I consider some of the most influential ‘Is’–‘Ought’ derivations. These include proposals by Prior and Searle. I argue that each proposal has a premise whose status as ethical or non-ethical is difficult to determine. In the second part of the paper, I consider taxonomies proposed by Karmo and Maitzen. I argue against both taxonomies. I end with the claim that we need a taxonomy of ethical sentences and that none of the current proposals are adequate.

Keywords ‘Is’–‘ought’ derivation · Taxonomy of ethical sentences · Ethical nihilism · Possible world · Ethical standard

In every system of morality, which I have hitherto met with, I have always remark’d, that the author proceeds for some time in the ordinary way of reasoning, and establishes the being of a God, or makes observations concerning human affairs; when of a sudden I am surpriz’d to find, that instead of the usual copulations of propositions, is, and is not, I meet with no proposition that is not connected with an ought, or an ought not. This change is imperceptible; but is, however, of the last consequence. For as this ought, or ought not, expresses some new relation or affirmation, ’tis necessary that it shou’d be observ’d and explain’d; and at the same time that a reason should be

S. Hill (✉)

Department of Philosophy, University of Massachusetts, 352 Bartlett Hall, Amherst, MA 01003, USA
e-mail: sahil0@philos.umass.edu

given, for what seems altogether inconceivable, how this new relation can be a deduction from others, which are entirely different from it.

This passage has spawned a lot of literature. Hume ([1739] 1992, p.469) seems to claim that there does not exist a valid argument that has all non-ethical sentences as premises and an ethical sentence as its conclusion. Starting with Prior, a number of counterexamples to this claim have been proposed. Unfortunately, all of these proposals are controversial. Even the most plausible have a premise that seems like it might be an ethical sentence or a conclusion that seems like it might be non-ethical. Since it is difficult to tell whether any of these counterexamples are genuine, we need a taxonomy that sorts out ethical sentences from non-ethical ones. We need to know the difference between an 'Is' and an 'Ought'. As we will see, though, it is not clear how to determine what that difference is.

In the first part of the paper, I establish the need for a taxonomy. I consider some of the most influential 'Is'–'Ought' derivations. These include proposals by Prior (1976) and Searle (1964). I argue that each proposal has a premise whose status as ethical or non-ethical is difficult to determine. In the second part of the paper, I consider taxonomies proposed by Karmo (1988) and Maitzen (1998). I argue against both taxonomies. I end with the claim that we need a taxonomy of ethical sentences and that none of the current proposals are adequate.

Some Influential 'Is'–'Ought' Derivations

Prior's proposal (1976) is a dilemma concerning the status of sentences that are the disjunction of a non-ethical sentence and an ethical sentence. Consider the sentences 'Tea-drinking is common in England' and 'All New Zealanders ought to be shot'. The first is non-ethical. The second is ethical. Now consider the sentence 'Either tea-drinking is common in England or all New Zealanders ought to be shot'. Either this sentence is ethical or it is non-ethical. Suppose it is ethical. Then you have the sound 'Is'–'Ought' derivation:

- (1) Tea-drinking is common in England.
- (2) Therefore, either tea-drinking is common in England or all New Zealanders ought to be shot.

Suppose instead that (2) is non-ethical. Then you have the valid 'Is'–'Ought' derivation:

- (1) It is not the case that tea-drinking is common in England.
- (2) Either tea-drinking is common in England or all New Zealanders ought to be shot.
- (3) Therefore, all New Zealanders ought to be shot.

So either way, whether (2) is ethical or non-ethical, you have got a valid 'Is'–'Ought' derivation. Here is the way Prior (1976, p. 91) puts it:

With regard to disjunctions of ethical and non-ethical statements, the 'autonomists' are in fact in the following dilemma: Either such statements are as a whole ethical or they are not. If they are, we may deduce ethical from non-

ethical propositions as [in the first argument]. And if they are not, we may deduce ethical from non-ethical propositions [as in the second argument].

It seems to me that Prior's proposal depends on the claim that if (2) has one classification in one of these arguments, then (2) must have the same classification in the other argument as well. If (2) is ethical in both arguments, then the first is sound. If (2) is non-ethical in both arguments, then the second is valid. In other words, Prior's proposal seems to depend on the claim that a sentence's classification as ethical or non-ethical is invariant and context independent. But there is another possibility. Perhaps (2) is non-ethical in the first argument and ethical in the second argument. (2) might seem to be non-ethical in the first argument because (2)'s truth is guaranteed by the truth of its non-ethical disjunct 'Tea-drinking is common in England'. On the other hand, (2) might seem to be ethical in the second argument because whether (2) is true depends on whether (2)'s ethical disjunct 'All New Zealanders ought to be shot' is true. Perhaps, in other words, a sentence's classification is sometimes variable and context dependent. If this were the case, then neither argument would be a successful derivation of an 'Ought' from an 'Is'. We need a taxonomy to see if this suggestion can be made to work.

Searle's proposal (1964) is a derivation of an ethical sentence from seemingly non-ethical sentences about utterances and promising. In many cases, a person can make a promise just by uttering a sentence. Typically, a person who has made a promise ought to keep it. Perhaps there are resources here for a successful 'Is'–'Ought' derivation. Maybe there is a way to derive an ethical conclusion about what someone ought to do from a sentence about what that person uttered without adding any ethical sentences along the way. This is Searle's proposal.

Here is the plan for this section: After stating Searle's derivation and the principles he employs, I formulate a definition for the derivation's key technical term, 'other things are equal'. I show that on one reading of this phrase, Searle's proposal has a premise that seems like it might be an ethical sentence. On another reading of the term, at least one of the principles behind Searle's derivation is false. I conclude with some remarks about how these considerations motivate the introduction of a taxonomy.

Searle's Argument

- (1) Jones uttered the words "I hereby promise to pay you, Smith, five dollars."
- (2) Conditions *C* obtain.
- (3) Jones promised to pay Smith five dollars.
- (4) Jones placed himself under an obligation to pay Smith five dollars.
- (5) Other things are equal.
- (6) Jones is under an obligation to pay Smith five dollars.
- (7) Jones ought to pay Smith five dollars.

Searle's proposal employs four principles. The first allows him to make the inference from (1) and (2) to (3). The conditions, *C*, that Searle talks about in (2) are supposed to be the set of necessary and sufficient conditions that must obtain in order for an utterance such as Jones' to count as a promise. Among the members of

C are the requirements¹ that the speaker must not be “under the influence of drugs” or “hypnotized” or “reporting an event” or “telling a joke” or “acting in a play” (1964, p. 45). Here is the first principle:

P1 If conditions C obtain, then for any subject, S , that is such that S utters “I hereby promise to perform action A ”, S promises to perform A .

The second principle allows Searle to make the inference from (3) to (4). According to this principle, placing oneself under an obligation is part of the “definition” and “analysis of the concept of promising” (1964, p.45–6). Here is the second principle:

P2 If a subject, S , promises to perform some action, A , then S places S under an obligation to perform A .

The third principle allows Searle to make the inference from (4) and (5) to (6). There is a set of necessary and sufficient conditions, C^* , that must obtain in order for an act of placing oneself under an obligation such as Jones’ to make one such that one is under an obligation. Among the members of C^* is the requirement² that Smith cannot say to Jones “I release you from your obligation” (1964, p.46). Here is the third principle:

P3 If other things are equal, then for any subject, S , that is such that S places S under an obligation to perform some action, A , S is obligated to perform A .

The fourth principle allows Searle to make the inference from (5) and (6) to (7). There is a set of necessary and sufficient conditions, C^{**} , that must obtain in order for one who, like Jones, is under an obligation to be such that one ought to do what one is under an obligation to do. Among the members of C^{**} is the requirement³ that Jones’ “obligation is [not] overridden by some other considerations, e.g., a prior obligation” (1964, p. 47). Here is the fourth and final principle.

P4 If other things are equal, then for any subject, S , that is such that S is under an obligation to perform some action, A , S ought to perform A .

Other Things Are Equal

Two of Searle’s principles, P3 and P4, require that “other things are equal”. So we must determine how Searle means to be using the *ceteris paribus* condition. Here is what he (1964, p. 47–8) says about the matter:

The force of the expression “other things being equal” in the present instance is roughly this. Unless we have some reason (that is unless we are *actually* prepared to give some reason) for supposing the obligation void... or the agent

¹ However, Searle says, “[t]his list will no doubt be somewhat indefinite because the boundaries of the concept of a promise, like the boundaries of most concepts in a natural language, are a bit loose” (1964, p. 45).

² As Searle (1964, p.46) points out, “it is possible for all sorts of things to happen which will release one from obligations one has undertaken.” So C^* has other members as well.

³ As Searle points out, there are other conditions that rule out cases in which “something extraneous to the relation of “obligation” to “ought”... interfere[s].”

ought not to keep the promise..., then the obligation holds and he ought to keep the promise.... It is sufficient to satisfy the condition that no reason to the contrary can *in fact* be given.... But unless *we* have some reason to the contrary, the *ceteris paribus* condition is satisfied, no evaluation is necessary, and the question whether he ought to do it is settled by saying “he promised.”

Searle (1964, p. 47) also says:

It is not part of the force of the phrase “other things being equal” that in order to satisfy it we need to establish that no reason *could ever be given by anyone* for supposing the agent is not under an obligation or ought not to keep the promise. That would be impossible and would render the phrase useless.

Searle considers two formulations of “other things are equal” in these passages. Both formulations concern whether or not some subject has a reason to suppose that Jones need not pay. According to the first passage, other things are equal “[u]nless... *we* are *actually* prepared” to give a reason otherwise and unless a reason “to the contrary can *in fact* be given.” So it seems like the reason in question has to be given by an actual subject. I assume that there is some connection⁴ between the reason given and the things prohibited by C^* and C^{**} . Perhaps it is that the reason must be one of the things prohibited by C^* or C^{**} . The second passage, on the other hand, says that “no reason *could ever be given by anyone* for supposing” Jones need not pay. So it seems that the reason in question can be given by any possible subject not just actual ones. Searle rejects this formulation. In order to keep these two formulations straight, call the second formulation “other things are equal*.” Given these passages, I take “other things are equal” to mean this⁵:

Other things are equal with respect to Jones’ case if and only if there does not exist an *actual* subject that either (i) judges some feature of Jones’ case, that is prohibited by C^* , to be a reason to suppose that Jones is not under an obligation or (ii) the subject in question judges some feature, that is prohibited by C^{**} , to be a reason to suppose that Jones ought not to do what he is under an obligation to do.

I take “other things are equal*” to mean this:

Other things are equal* with respect to Jones’ case if and only if there does not exist a *possible* subject that either (i) judges some feature of Jones’ case, that is prohibited by C^* , to be a reason to suppose that Jones is not under an obligation or (ii) the subject in question judges some feature, that is prohibited by C^{**} , to be a reason to suppose that Jones ought not to do what he is under an obligation to do.

⁴ I must admit that Searle does not explicitly say this. So I am being a little loose with the text here. But I have a justification for this attribution. As we will see later, it is needed to explain why Searle endorses the formulation of ‘other things are equal’ in the first passage and resists the formulation of ‘other things are equal’ in the second passage. Without stipulating some kind of connection like this, Searle’s preference cannot be explained.

⁵ So the rough idea is: other things are equal (or other things are equal*) if and only if it is not the case that (i) there is some actual person (or possible person) who judges that ‘Jones ought to pay’ is false and (ii) ‘Jones ought to pay’ is false. I think this is the core idea. But it does not capture all the details.

Either Searle's Argument is Invalid or (5) is Ethical

In the remainder of this section, my strategy will be as follows: First, I consider the formulation of the *ceteris paribus* condition that Searle rejects and argue that given this formulation (5) seems like it might be an ethical sentence. Next, I consider the formulation of "other things are equal" that Searle endorses and argue that given this formulation either P3 or P4 is false. I take this to show that the formulation according to which (5) seems like an ethical sentence is preferable and therefore we need a taxonomy of ethical sentences to tell us how to classify (5).

(5) as "Other Things are Equal"

As we have seen, Searle takes pains to say that it "is not part of the force of the phrase" that other things are equal* and that this somehow "justif[ies] [his] claim that the satisfaction of [the *ceteris paribus*] condition does not necessarily involve anything evaluative." The concern seems to be that if (5) is read as "other things are equal**", then (5) is an ethical sentence.

There is an argument that gives voice to Searle's concern. You can think of the argument as having four steps. In step 1, I restate the definition of other things are equal* and I point out that there exists a possible agent that is ideal. In step 2, I use step 1 to show that other things are equal* *only if* 'Jones ought to pay' is true. In step 3, I use step 1 to show that other things are equal* *if* 'Jones ought to pay' is true. In step 4, I put steps 2 and 3 together and conclude that other things are equal* *if and only if* 'Jones ought to pay' is true. I point out that this makes (5) look like an ethical sentence.

- Step 1 Other things are equal* with respect to the case of Jones' utterance if and only if there does not exist a feature of that case that is such that either (i) (a) there exists a possible subject that judges 'Jones is obligated to pay' is false and (b) some member of C* fails to obtain or (ii) (a) that possible subject judges 'Jones ought to pay' to be false and (b) some member of C** fails to obtain. Now, there exists a possible subject that is an ideal agent. Being ideal, she is maximally opinionated. For any case whatsoever, she has got an opinion about whether or not someone in that case is obligated or ought to do something. Being ideal, she is also perfectly reliable. For any case you give her, the ideal agent's judgments about whether someone in that case is obligated or ought to do something will be completely in step with whether the members of C* and C** obtain in that case. If one of the members of one of these sets does not obtain in that case, then she judges that it is not the case that the person is obligated or ought to do something.
- Step 2 Now, since she is among the possible agents, other things are equal* with respect to Jones' case only if it is not the case that there is some feature of the case such that (i) (a) the ideal agent judges to be a reason to suppose 'Jones is obligated to pay' is false and (b) some member of C* fails to obtain or (ii) (a) the ideal agent judges that feature to be a reason to suppose 'Jones ought to pay' is false and (b) one of the members of C** does not obtain. So both (i) and (ii) have got to be false in order for other things to be equal*. Since she is maximally opinionated, the ideal agent has a judgment about whether

some feature of Jones' case is a reason to suppose either 'Jones is obligated to pay' or 'Jones ought to pay' is false. Since she is perfectly reliable, if C^* or C^{**} is not satisfied, then she makes the corresponding judgment that some feature of Jones' case is a reason to suppose that either 'Jones is obligated to pay' or 'Jones ought to pay' is false. In other words, if (ii) is satisfied, then so is (i). So other things are equal* only if it is not the case that one of the members of C^* or C^{**} fails to obtain. In other words, other things are equal* only if 'Jones is obligated to pay' and 'Jones ought to pay' are true.

- Step 3 Again, because she is maximally opinionated, we have a guarantee that the ideal agent makes a judgment about Jones' case. She has got an opinion about whether Jones is obligated to pay and an opinion about whether he ought to pay. What is more, because she is perfectly reliable, we have a guarantee that the ideal agent's judgment is in step with whether C^* and C^{**} obtain in Jones' case. So if some member of C^* fails to obtain, then she judges that 'Jones is obligated to pay' is false. If some member of C^{**} fails to obtain, then she judges that 'Jones ought to pay' is false. Thus, if either C^* or C^{**} fail to obtain, then either (1) (a) or (2) (a), the conditions about the judgment of a possible agent, and either (1) (b) or (2) (b), the conditions about C^* and C^{**} failing to obtain, are not satisfied. So other things are equal*, if it is not the case that one of the members of C^* or C^{**} fails to obtain. In other words, other things are equal*, if 'Jones is obligated to pay' and 'Jones ought to pay' are true.
- Step 4 Putting together the results of steps 2 and 3 we get: other things are equal* if and only if 'Jones is obligated to pay' and 'Jones ought to pay' are true. But this makes it seem like (5) is an ethical sentence. Or, anyway, the truth conditions of (5) are identical to the truth conditions of the conjunction of the derivation's conclusion, (7), with 'Jones is obligated to pay'. So if you go with this reading, Searle's argument could have been a lot shorter. He could have just said: Jones is obligated to pay and Jones ought to pay. Therefore, Jones ought to pay. For this reason it seems like Searle had better go with the other reading of "other things are equal". He was right to resist reading it as other things are equal*⁶.

⁶ My interpretation explains why Searle (1964, p. 47) thinks that reading "other things are equal" as other things are equal "justif[ies] [his] claim that the satisfaction of [the *ceteris paribus*] condition does not necessarily involve anything evaluative." Recall that the big difference between this way of reading the phrase and reading it as other things are equal* has to do with whether or not you allow possible agents in as candidates for satisfying this requirement. An essential feature of the argument that (5) is ethical was my appeal to the existence of a possible ideal agent. Now, it is not too controversial to suppose that there is a possible agent that is ideal. Using this supposition we could show that (5) looks like it might be ethical. It is much more controversial, however, to say that some actual agent is ideal. For this reason, a similar argument cannot be used to show that (5), on Searle's preferred reading, is ethical. So my interpretation is able to explain why Searle thinks that reading (5) as saying that other things are equal makes (5) a non-ethical sentence.

(5) as “Other Things are Equal”

We can now devise a test for Searle’s derivation. Searle says that in order for other things to be equal, it is not necessary that other things are equal*. So it is acceptable if some possible agent judges that some feature of Jones’ case makes him such that he ought to pay as long as no actual subject makes such a judgment. So suppose that other things are equal. Suppose also that it is not the case that other things are equal*.

Since other things are equal, there does not exist a feature of Jones’ case that is such that (i) (a) an actual subject judges ‘Jones is obligated to pay’ is false and (b) that feature obtains if and only if C^* fails to obtain or (ii) (a) the actual subject judges that ‘Jones ought to pay’ is false and (b) that feature obtains if and only if some member of C^{**} does not obtain. According to P3, if other things are equal and Jones places himself under an obligation to pay, then Jones is under an obligation to pay. So P3 says that Jones is under an obligation to pay. According to P4, if other things are equal and Jones is under an obligation to pay, then Jones ought to pay. So P4 says that Jones ought to pay.

Since it is not the case that other things are equal*, there exists some feature of Jones’ case that is such that (i) (a) a possible subject judges that some feature of Jones case is a reason to think that ‘Jones is obligated to pay’ is true and (b) that feature is prohibited by a member of C^* or (ii) (a) that possible subject judges that some feature of Jones’ case to be a reason to think that ‘Jones ought to pay’ is false and (b) that feature is prohibited by a member of C^{**} .

We can show that either P3 or P4 is false. Since other things are not equal* either some member of C^* or C^{**} does not obtain. But if a member of C^* does not obtain, then ‘Jones is obligated to pay’ is false. Since P3 says that ‘Jones is obligated to pay’ is true, P3 is false. If instead a member of C^{**} does not obtain, then ‘Jones ought to pay’ is false. But P4 says that ‘Jones ought to pay’ is true. So P4 is false. Either way one of Searle’s principles, on this reading, is going to be false.

Since the argument is invalid if “other things are equal” is read in this way, it seems like our best bet is to go with the formulation of the *ceteris paribus* condition that Searle rejects—other things are equal*. The worry about this way of reading Searle was that (5) seemed like it might be an ethical sentence. (5)’s status is controversial. We need a taxonomy of ethical sentences to decide this.

Taxonomies

The interesting proposals considered above each have a premise whose status as ethical or non-ethical is difficult to determine. Below I consider taxonomies proposed by Karmo and Maitzen and the classifications these taxonomies give for the sentences in dispute.

Consider Karmo’s taxonomy (1988). Karmo analyzes ethicalness in terms of possible worlds and ethical standards. So I need to say a bit about what he means by each of these.

Possible Worlds Some sentences are uncontroversially non-ethical. ‘Tea-drinking is common in England’ is one of these. A sentence is uncontroversially non-ethical if

and only if every member of “the logical-autonomy-of-ethics⁷ debate would unite in calling [it] non-ethical” (Karmo 1988, p. 254). There is a set of all uncontroversially non-ethical sentences. You can pick out a possible world by assigning truth-values to all members of this set. There is one possible world that is actual.

Ethical Standards Some sentences are uncontroversially ethical. ‘It is wrong to purchase coffee that is not fairly traded’ is one of these. A “sentence is ‘uncontroversially ethical’ just in case all parties to the logical-autonomy-of-ethics debate would unite in calling it ethical” (Karmo 1988, p. 254). There is a set of all uncontroversially ethical sentences. You can determine an ethical standard by assigning truth-values to all members of this set. There is one ethical standard that is correct.

Talk about the truth-value of a sentence will always have a form like this: ‘Sentence, *S*, is true (or false) at a possible world, *W*, with respect to an ethical standard, *E*’. Here is Karmo’s taxonomy:

A sentence, *S*, is ethical at a possible world, *W*, if and only if (i) *S* is true at *W* with respect to one ethical standard, *E*, and (ii) *S* is false at *W* with respect to another ethical standard, *E**.

Karmo’s taxonomy has the result that a sentence may be ethical under some conditions and non-ethical under other conditions⁸. Call a taxonomy like this a variable context-dependent taxonomy⁹. An interesting feature of Karmo’s taxonomy is that, if it is correct, there exists a proof showing that there are no sound ‘Is’–‘Ought’ derivations—although there are some valid ones¹⁰.

Next consider Maitzen’s taxonomy (1998). For Maitzen, a sentence’s status as ethical is linked to the instantiation of moral properties. *S*, he says, “is an ethical sentence only if it is impossible for [*S*], standardly construed, to be true unless a given moral property (such as moral rightness or wrongness) is possessed by at least one object” (Maitzen 1998, p. 361). Notice that Maitzen gives only a necessary condition that a sentence needs to satisfy in order to count as ethical. It is not sufficient.

⁷ The logical-autonomy-of-ethics debate is the debate between those who believe you can derive an ‘ought’ from an ‘is’ and those who believe that you cannot.

⁸ As Karmo points out, this idea appears in Humberstone (1982). It also appears in Humberstone (1996).

⁹ For other variable and context-dependent approaches see Pigden (1989) and Schurz (1994) and (1997). Thanks to a referee for this journal for helping me to see the importance of making the context dependent/independent distinction explicit.

¹⁰ Here is Karmo’s proof: “if sentences *S*₁, ..., *S*_{*n*} (where *n*>0) entail sentence *S*(*n*+1), then for any possible world *W* in which *S*(*n*+1) is ethical, if all of *S*₁, ..., *S*_{*n*} are true in *w*, then at least one of *S*₁, ..., *S*_{*n*} is ethical in *w*. (Proof: since all of *S*₁, ..., *S*_{*n*} are true in *w*—that is, are true in *w* with respect to the correct ethical standard *E*—*S*(*n*+1) is true in *w* with respect to *E*. If *S*(*n*+1) is to be ethical in *w*, then there must be some alternative ethical standard *E** with respect to which *S*(*n*+1) is false in *w*. Since replacing *E* with *E** changes the truth value of *S*(*n*+1) in *w*, and *S*(*n*+1) is entailed by *S*₁, ..., *S*_{*n*}, replacing *E* with *E** changes the truth value in *w* of at least one of *S*₁, ..., *S*_{*n*}; so at least one of *S*₁, ..., *S*_{*n*} is ethical in *w*.” (p 256–7).

I need to say a bit about what Maitzen means by the phrase “a given moral property”. According to him (1998, p. 362), the sentence ‘Some ethical sentences are true’ is not such that if it is true, then a given moral property is instantiated. Consider two possible worlds. At one world rightness is the only moral property instantiated and the sentence ‘Something is right’ is true. At the other world wrongness is the only moral property instantiated and the sentence ‘Something is wrong’ is true. At both worlds the sentence ‘Some ethical sentences are true’ is true. At both worlds “some moral property or other” is instantiated. But at one of the worlds rightness is all that is instantiated and at the other world wrongness is all that is instantiated. So “a given moral property” is not instantiated. Given moral property talk, then, amounts to this: In order for a sentence to be classified as ethical, there has got to be some moral property that is instantiated at every world at which that sentence is true. Here is Maitzen’s taxonomy:

A sentence, *S*, is ethical only if there exists a moral property, *P*, that is such that for any possible world, *W*, at which *S* is true, *P* is instantiated at *W*.

Maitzen’s taxonomy has the result that a sentence is either ethical under all conditions or ethical under none. Call a taxonomy like this an invariant and context-independent taxonomy¹¹. An interesting feature of Maitzen’s taxonomy is that if it is correct, then there are sound ‘Is’–‘Ought’ derivations. I will say more about this below.

The Bearing of these Taxonomies on Prior and Searle’s Proposals

Now consider how these initially plausible taxonomies bear on Prior and Searle’s proposals.

Prior

The controversial sentence in Prior’s proposal is (2). This sentence read ‘Either tea-drinking is common in England or all New Zealanders ought to be shot’. Maitzen classifies (2) as non-ethical. (2) is true at any possible world at which the English drink a lot of tea. There is no moral property *P*, such that *P* is instantiated at every world at which (2) is true¹². So, according to Maitzen’s taxonomy, (2) is non-ethical.

Karmo classifies (2) as ethical in some worlds and non-ethical in other worlds. Pick out any possible world, say *W*, at which ‘Tea-drinking is common in England’

¹¹ Recall Prior’s example. If you agree that the sentence ‘Either tea-drinking is common in England or all New-Zealanders ought to be shot’ is either ethical in all contexts or non-ethical in all contexts, then either way you get a valid ‘Is’–‘Ought’ derivation.

¹² You may be worried that if tea drinking is common in England, then some moral property or other is bound to be instantiated. That is probably right. But that does not mean that some given moral property is instantiated. I think, for Maitzen, the idea is something like this: Pick any moral property you want, say wrongness. There is a possible world where nobody does anything wrong. The English all treat each other very well and so wrongness is not instantiated at that world. So a given moral property is not instantiated.

is true. (2) is non-ethical at W . To see this, determine any ethical standard that you want. The truth-value of ‘Tea-drinking is common in England’ does not vary with respect to different ethical standards. So no matter what truth-values you assign to the set of all uncontroversially ethical sentences, (2) is going to be true at W because (2)’s non-ethical disjunct is true at W .

Now pick out any possible world, say W^* , at which ‘Tea-drinking is common in England’ is false. Since the English do not drink much tea at W^* , whether (2) is true depends on whether its ethical disjunct ‘All New Zealanders ought to be shot’ is true. So determine two ethical standards, E and E^* . For E , assign ‘true’ to ‘All New Zealanders ought to be shot’¹³. For E^* , assign ‘false’ to this sentence. (2) is true at W^* with respect to E and false at W^* with respect to E^* . Karmo’s taxonomy classifies (2) as ethical at W^* . Recall that Prior’s proposal consists of two arguments. First there is:

- (1) Tea-drinking is common in England.
- (2) Therefore, either tea-drinking is common in England or all New Zealanders ought to be shot.

The second argument is:

- (1) It is not the case that tea-drinking is common in England.
- (2) Either tea-drinking is common in England or all New Zealanders ought to be shot.
- (3) Therefore, all New Zealanders ought to be shot.

The two taxonomies give different results: According to Maitzen’s invariant and context-independent taxonomy, (2) is non-ethical in both arguments. So the first argument fails and the second succeeds. According to Karmo’s variant and context dependent taxonomy, (2) is non-ethical in the first argument and ethical in the second. So neither argument is successful.

Searle

The controversial premise in Searle’s proposal is (5). It read ‘other things are equal*’. I argued that (5) is true if and only if the sentences ‘Jones is obligated to pay’ and ‘Jones ought to pay’ are both true.

According to Maitzen’s taxonomy, (5) is ethical. The truth conditions of (5) are identical to the truth conditions of ‘Jones is obligated to pay and Jones ought to pay’. So at any world where ‘Jones ought to pay’ is false, (5) is also false. Now, there does not exist a possible world at which ‘Jones ought to pay’ is true and Jones fails to instantiate the moral property of ought-to-pay-Smith-five-dollars-ness. So for Maitzen, (5) is ethical¹⁴.

¹³ Although there maybe be an ethical standard according to which all New Zealanders ought to be shot, such a horrible standard is, of course, not the correct one.

¹⁴ Well I’m exaggerating a little bit. Maitzen’s taxonomy is only a partial one. His taxonomy tells you when a sentence is definitely not ethical. But it does not tell you when a sentence definitely is ethical.

Karmo's taxonomy agrees with Maitzen's. Determine two ethical standards E and E^* . 'Jones is obligated to pay' and 'Jones ought to pay' are both true with respect to E and both are false with respect to E^* . Now, the truth conditions of (5) are identical to the truth conditions of 'Jones is obligated to pay and Jones ought to pay'. Thus, for any possible world you pick, say W , (5) is true with respect to E and false with respect to E^* . So for Karmo, (5) is ethical.

Maitzen's Criticism of Karmo

In this Section I consider Maitzen's criticism of Karmo¹⁵. Maitzen argues that Karmo's taxonomy misclassifies non-ethical sentences as ethical. It seems to me that his criticism can be understood in two ways. Understood in one way, Maitzen argues that the sentence 'Either no ethical sentence is true or torturing babies is wrong' is a non-ethical sentence. But since Karmo's taxonomy classifies this sentence as ethical, his taxonomy is incorrect. Understood in another way, Maitzen appeals to the intuition that Ethical Nihilism is a coherent theory. But since Karmo cannot accommodate this intuition, his taxonomy is incorrect. I consider both understandings of Maitzen's complaint and argue that Karmo need not worry about either of them.

The Possibility of Ethical Nihilism as a Premise

Consider the first way of understanding Maitzen's criticism: Karmo's taxonomy classifies the sentence 'Either no ethical sentence is true or torturing babies is wrong' as ethical at all possible worlds¹⁶. Maitzen (1998, p. 354–6) sets out to argue against

¹⁵ For another interesting objection to Karmo's taxonomy see Nelson (2007). He argues against the claim that Karmo's taxonomy has the result that there is no sound derivation of an 'ought' from an 'is'. Consider this argument:

- (1) Stalin authorized the Katyn Massacre.
- (2) Therefore, it is not the case that both 'If Stalin authorized the Katyn Massacre, then Stalin was evil' and 'It is not the case that Stalin was evil'.

Nelson argues that Karmo's taxonomy classifies (1) as non-ethical and (2) as ethical. This is a topic for another paper and there is not space to go into it here. But it seems to me that Karmo's taxonomy classifies (2) as non-ethical at every world at which (1) is true. So Karmo's claim that his taxonomy gives the result that there are no sound 'Is'–'Ought' derivations goes unscathed.

¹⁶ To see this, consider two ethical standards, say E and E^* . With respect to E , the sentences 'No ethical sentence is true' and 'Torturing babies is wrong' are both false. With respect to E^* , the sentence 'No ethical sentence is true' is false and the sentence 'Torturing babies is wrong' is true. Now pick any possible world you want, say W . At W , the sentence 'Either no ethical sentence is true or torturing babies is wrong' is false with respect to E and true with respect to E^* . So the sentence in question is ethical at W . But W can be any possible world. So Maitzen is right to claim that, for Karmo, the sentence 'Either no ethical sentence is true or torturing babies is wrong' is ethical at all possible worlds.

Karmo's classification and places himself under the constraint stated in the passage below:

[Karmo's] taxonomy fails to reckon with a certain kind of sentence which, *given non-question-begging assumptions*, can be shown to be, when standardly construed, non-ethical in at least one world even though it satisfies his conditions for being an ethical sentence in every world.... I recognize the contentiousness of [my assumption]... but I hasten to point out that *the assumption begs no relevant questions* in my dispute with Karmo.

In light of this, it seems to me that each premise of Maitzen's argument must have a formulation in terms of Karmo's taxonomy that is true. If not, then Maitzen's argument begs the question against Karmo.

Maitzen's Argument

The first premise of Maitzen's argument is what Karmo should challenge. It is this:

(M) Possibly Ethical Nihilism

Now, (M) needs an intelligible formulation in terms of Karmo's taxonomy. Remember that talk about sentences, for Karmo, is always relativized to possible worlds and ethical standards. What is more, Maitzen (1998, p. 355) says that Ethical Nihilism (from now on EN) is the proposition "that no ethical sentence, standardly construed, is true." I can therefore think of only four ways to construe (M) so that it is intelligible:

- (M₁) EN is true at the actual world with respect to the correct ethical standard.
- (M₂) EN is true at the actual world with respect to some ethical standard, say *E*.
- (M₃) EN is true at some possible world, say *W*, with respect to the correct ethical standard.
- (M₄) EN is true at some possible world, say *W*, with respect to some ethical standard, say *E*.

So (M) is intelligible. In order to satisfy Maitzen's promise not to beg any questions, at least one of (M)'s formulations must be true. There is an argument that shows that each formulation is false¹⁷. There are two points to develop before I give the argument.

¹⁷ Maitzen (1998, p. 357) anticipates this argument and offers a response to it:

Surely, moreover, there is such a thing as the 'nihilistic ethical standard,' the standard that assigns 'false' to every member of the set U of uncontroversially ethical sentences.... One might reply that some pairs of ethical sentences... are logical contradictories.... [O]f course ['It is wrong to purchase coffee that is not fairly traded'] and ['It is not the case that it is wrong to purchase coffee that is not fairly traded'] are contradictories, but ['It is not the case that it is wrong to purchase coffee that is not fairly traded'] is not then an ethical sentence, for reasons that will emerge shortly.

Maitzen's formulation of the argument depends on the premise that \neg FT is uncontroversially ethical with respect to Karmo's taxonomy. Since Maitzen doesn't classify \neg FT as ethical, I concede that it is not. But my version of the argument depends only on the premise that, as I show in the main text, \neg FT is ethical with respect to Karmo's taxonomy even though it is not uncontroversially ethical.

First, recall the distinction, for Karmo, between ethical sentences and uncontroversially ethical sentences. Consider the sentence “It is wrong to purchase coffee that has not been fairly traded”. Call this sentence FT. FT is uncontroversially ethical. Now consider the sentence “It is not the case that it is wrong to purchase coffee that has not been fairly traded”. Call this sentence \neg FT. \neg FT is not uncontroversially ethical. For according to Karmo (1988, p. 254), only sentences such that “all parties to the” debate in question “unite in calling... ethical” count as uncontroversially ethical. Now, \neg FT just says that purchasing unfairly traded coffee lacks the moral property wrongness. So \neg FT can be true at a world at which wrongness is never instantiated. But according to Maitzen’s taxonomy an ethical sentence must, among other things, be such that there is a moral property that is instantiated at every world at which that sentence is true. So Maitzen does not unite with Karmo in calling \neg FT ethical. Therefore, \neg FT is not uncontroversially ethical. This completes the first point.

The second point to develop is this: \neg FT is ethical by Karmo’s standards even though it is not uncontroversially ethical. To see this reason as follows: Both Karmo and Maitzen agree that FT is ethical. So it is safe to say that FT is uncontroversially ethical. Now pick any possible world, say W , and determine any ethical standard you want, say E . Since an ethical standard is any complete assignment of truth-values to the set of all uncontroversially ethical sentences, you will have to assign either ‘true’ or ‘false’ to FT. Suppose that in determining E you assign ‘true’ to FT. Then \neg FT will be false at W with respect to E . Now determine another ethical standard, call it E^* , by assigning ‘false’ to FT. Then \neg FT is true at W with respect to E^* . So the truth-value of \neg FT varies at W depending on whether the ethical standard you determine is E or E^* . But this sort of variation is exactly what makes a sentence ethical for Karmo. So although it is not uncontroversially ethical, \neg FT is ethical. This completes the second point.

In light of these two points we are now in a position to give an argument against (M): Pick any possible world you want, say W . Now determine an ethical standard, call it E , by assigning ‘false’ to all members of the set of uncontroversially ethical sentences. Since it is uncontroversially ethical, FT is false at W with respect to E . Since FT is false at W with respect to E , \neg FT is true at W with respect to E . Since \neg FT is both an ethical sentence and true, the sentence ‘No ethical sentence is true’ is false at W with respect to E . Now determine any other ethical standard you want, say E^* . Just make sure E^* is distinct from E . However you decide to determine E^* , you are going to have to assign ‘true’ to at least one member of the set of uncontroversially ethical sentences. If you do not do that, and you assign ‘false’ to every member, then you will just end up determining E all over again. Since uncontroversially ethical sentences are also ethical sentences, the sentence ‘No ethical sentence is true’ is false at W with respect to E^* . Thus, for any ethical standard you determine and for any possible world you pick, the sentence ‘No ethical sentence is true’ is false. Therefore, (M₁) and (M₂) and (M₃) and (M₄) are all false. So (M) is false.

The Possibility of Ethical Nihilism as a Counterexample

There is another way to read Maitzen. Perhaps Maitzen is simply appealing to the intuition that Ethical Nihilism is not inconsistent. Karmo's taxonomy does not accommodate this intuition. So it is incorrect. In some places Maitzen (1998) argues as though this is what he is doing. What is more, in another paper, and in a different context not concerning Karmo, Maitzen (2006, p456) makes a point that can be used to motivate this line. He says:

[I]f [EN] is an ethical proposition, it is in effect self-refuting. The wide variety of ethical nihilists in the recent history of philosophy would rightly reject this too-quick, taxonomic refutation of their position. Assuming that ethical nihilism is not trivially false, [EN] cannot belong to the domain of ethics, and so neither can [¬EN]. The same lesson applies to any domain about which nihilism is not self refuting.

This is not persuasive. You can formulate a theory with Karmo's taxonomy that is nearly identical to Maitzen's version of Ethical Nihilism. Consider Maitzen's formulation (1998, p. 355). He stipulated that Ethical Nihilism is "the thesis that no ethical sentence, standardly construed, is true." Consider Maitzen's taxonomy. For him a sentence is ethical only if there exists a moral property that is such that for any possible world at which that sentence is true, the moral property in question is instantiated. Now combine Maitzen's formulation of Ethical Nihilism with Maitzen's taxonomy. Call it Maitzen's Ethical Nihilism (or MEN). MEN assigns 'true' to the sentence EN and MEN assigns 'false' to any sentence that Maitzen classifies as ethical.

Consider an alternative formulation of Ethical Nihilism. Let Ethical Nihilism be the thesis that any sentence that Maitzen's taxonomy classifies as ethical is false. Or put another way, consider every sentence such that there is a moral property that is instantiated at every possible world at which that sentence is true. Let ethical nihilism be the claim that all such sentences are false. Now combine this formulation with Karmo's taxonomy. Call it Karmo's Ethical Nihilism (or KEN). KEN assigns 'false' to EN and KEN also assigns 'false' to any sentence that Maitzen classifies as ethical.

It seems to me that EN is the only sentence about which MEN and KEN disagree. MEN says EN is true while KEN says EN is false. They both agree, however, that any sentence Maitzen classifies as ethical is false. Now if you really had the intuition that there has got to be some coherent way to formulate Maitzen's specific version of Ethical Nihilism down to every last detail and you thought that Karmo's taxonomy cannot accommodate this intuition, then you are right about that. But if instead your intuition was just that there must be some coherent way of formulating Ethical Nihilism somewhere in the neighborhood of Maitzen's version, but you don't really care about which truth-value it assigns to EN, then I can put your concerns to rest. Karmo's taxonomy can accommodate this intuition.

Here is another way to put the point: Ethical Nihilism should not be formulated as some metalinguistic claim about which sentences are true and which sentences are

false. Ethical Nihilism is just the idea that nothing is good or evil or right or wrong. Nothing has any moral properties¹⁸.

Criticism of Maitzen's Taxonomy

If my arguments are sound, then Karmo should not be bothered by either understanding of Maitzen's criticism. I now turn to what seem to me to be problems with Maitzen's taxonomy. Maitzen (1998) and (2006) offers the following proposal as an 'Is'–'Ought' derivation:

- (1) Some ethical sentences are true.
- (2) Either no ethical sentence is true or torturing babies is wrong.
- (3) Therefore, torturing babies is wrong.

We have already considered Maitzen's classification of (1) and (2). Concerning (1) he says (1998) that the sentence is non-ethical "since to say [(1)] is true is not to commit oneself to the instantiation of a *given* moral property (only, at most, to the instantiation of some moral property or other)." Concerning (2) he says (2006, p458) that "[I]n at least one possible world, [(2)] falls outside the domain of ethics. [What is more], a proposition's status as outside a domain is an essential feature of it: if there is any possible world in which [(2)] falls outside the domain of ethics, it does so in every possible world in which it exists." So the idea is that there exists a possible world where (2) is true and no moral properties are instantiated. So (2) is non-ethical.

It seems to me that this reveals two problems for Maitzen's taxonomy. First, the conclusion of his 'Is'–'Ought' derivation is 'Torturing babies is wrong'. But his taxonomy does not classify this sentence as ethical. To see this reason as follows: There exists some possible world, W , at which nobody ever tortures a baby. Well, just because nobody tortures a baby at W it does not follow that it is permissible for people at W to do so. It seems like 'Torturing babies is wrong' is true at W regardless of whether anybody does it. But since nobody ever tortures babies at W , wrongness is never instantiated by the act of torturing babies. Let us further say that W is a world at which nobody ever does anything wrong at all. So you end up with a world, W , that is such that (1) 'Torturing babies is wrong' is true at W and (2) wrongness is not instantiated at W . The result is that 'Torturing babies is wrong' is a non-ethical

¹⁸ The problem I have been discussing in this section of the paper is related to what Pigden (2007) calls the Doppelgänger Problem. The problem is this: Ethical Nihilism is the thesis that every ethical sentence is false. Suppose that S is an ethical sentence. If S is ethical, then $\neg S$ is ethical too. But either S or $\neg S$ has got to be true. So Ethical Nihilism is false. I propose that we understand Ethical Nihilism in terms of moral properties rather than in terms of which sentences are true and which sentences are false. Pigden deals with this problem differently. He continues to construe Ethical Nihilism in terms of sentences. Pigden's strategy is to reformulate Ethical Nihilism as a claim about a restricted set of ethical sentences, what he calls "non-negative atomic moral judgments." These sentences are such that S 's belonging to that restricted set does not entail $\neg S$'s also belonging to that set. Moreover, if S is a non-negative atomic moral judgment, then $\neg S$ is not. Formulating Ethical Nihilism as the claim that all non-negative atomic moral judgments are false, then, does not get you into a contradiction.

sentence. So if Maitzen's taxonomy is correct, the conclusion of his 'Is'–'Ought' derivation, is non-ethical¹⁹.

Second, if Maitzen's taxonomy is correct, then sound 'Is'–'Ought' derivations are far too easy to construct. Consider the following argument:

- (1) Either the Holocaust was good or the Holocaust was evil.
- (2) It is not the case that the Holocaust was good.
- (3) Therefore, the Holocaust was evil.

Maitzen's taxonomy classifies (1) as non-ethical. Granted, if (1) is true, then some moral property, either goodness or evilness, is instantiated. But the truth of (1) does not require that one or the other in particular is instantiated. So (1) is non-ethical. Maitzen's taxonomy also classifies (2) as non-ethical. As we have seen, he classifies ~FT as non-ethical. His reasoning was that there is a possible world at which ~FT is true and no moral property is instantiated. According to the same reasoning (2) is non-ethical. (2) might be true because the Holocaust was evil. But (2) can be true at a world at which evil is never instantiated. Perhaps the Holocaust never happened at that world. So (2) is non-ethical. What about (3)? The only way to make (3) true is to make evil instantiated by the Holocaust. So Maitzen classifies (3) as ethical. If his

¹⁹ A reviewer gives a response on behalf of Maitzen: Consider the universal being-an-act-of-torturing-babies. This universal exists at *W*. Consider the property is-a-wrong-making-feature. This is a moral property. What is more, this property is instantiated by being-an-act-of-torturing-babies at *W*. So it seems like for any world at which (3) is true, the moral property is-a-wrong-making-feature is instantiated. Therefore, Maitzen can classify (3) as an ethical sentence.

I think that this is a really interesting and promising way for Maitzen to respond. I do not have a full counter-response worked out. But let me give you a sketch of a counter-response. I suspect that just as it is difficult to tell, in some cases, whether a sentence is ethical or non-ethical, it will be difficult, in some cases, to tell whether a property is moral or non-moral. I suspect that different plausible answers to the question "What is a moral property?" will give you different answers to the question "Are some properties of universals such as is-a-wrong-making-feature moral properties?" In other words, I think that is—a wrong-making-feature's status as moral or non-moral is controversial. So we won't know whether my objection to Maitzen's taxonomy is successful until we know what counts as a moral property.

Now, Maitzen takes moral properties to be undefined and primitive. He says that "in order to avoid an infinite regress, we may have to take something, perhaps the notion of moral property, as primitive, as undefined, but then any (finite) conceptual analysis must take something as primitive" (Closing the 'Is'–'Ought' Gap, p 362). I have no complaints about taking some concepts to be primitive in an analysis. But I do not think that taking moral properties to be primitive is an option for Maitzen. Here is why: To my mind, the motivation for introducing taxonomies of ethical sentences in the first place is the need to evaluate alleged 'Is'–'Ought' derivations. There are sentences whose status as ethical or nonethical is difficult to determine. We need to know how to classify them. But Maitzen's proposal defines ethical sentences in terms of something that is equally problematic—moral properties. There are cases, such as the one the reviewer raises, of properties whose status as moral or non-moral is difficult to assess. Until we know what a moral property is, we cannot know whether Maitzen's taxonomy is adequate and, therefore, we cannot know which way to classify the controversial sentences appearing in Prior and Searle's 'Is'–'Ought' derivations. In other words, the same problem reemerges. So if Maitzen takes moral properties to be primitive, then the motivation for introducing his taxonomy—the ability to evaluate Prior and Searle's proposals—is undermined. So Maitzen owes us an account of what a moral property is.

In light of the reviewer's point, I would restate my objection as a challenge. If you can tell me what a moral property is, then I'll tell you whether Maitzen's taxonomy is correct. If you cannot tell me what a moral property is, then we had better look for another taxonomy that does not take moral properties to be primitive because otherwise the same problems about determining the status of ethical sentences remerge with determining the status of moral properties. Again, what I have here is just a sketch. I think the reviewer's comment opens up an interesting line of response for Maitzen. There is more to be said here.

taxonomy is correct there is a sound ‘Is’–‘Ought’ derivation that has as its conclusion one of the most substantial ethical sentences that you can imagine. This generalizes. If Maitzen’s taxonomy is correct, then you can validly derive any ethical sentence you want from non-ethical premises. Pick your favorite ethical sentence, say ‘*O* instantiates moral property *P*’. Now construct a sentence such as ‘Either *O* instantiates moral property *P* or *O* instantiates moral property *Q*’. Construct another sentence such as ‘It is not the case that *O* instantiates *Q*’. Both sentences are non-ethical. Use them as premises and you can derive the ethical sentence ‘*O* instantiates moral property *P*’. It seems to me that this is not plausible. It makes ‘Is’–‘Ought’ derivations far too easy to construct²⁰.

Criticism of Karmo’s Taxonomy

In the last Section I considered problems with Maitzen’s taxonomy. Now I turn to problems with Karmo’s taxonomy.

Karmo’s Taxonomy Misclassifies EN

Earlier, I argued that for any possible world and for any ethical standard, EN is false at that world and with respect to that ethical standard. Since EN is false at all worlds and with respect to all ethical standards, Karmo’s taxonomy classifies EN as non-ethical. But EN is an ethical sentence. So Karmo’s taxonomy is incorrect.

I can only justify the premise that EN is an ethical sentence with an appeal to intuition. I am not alone in having this intuition²¹. But if you believe otherwise, then I do not think I have the resources to convince you to come over to my side²². However, I think that I can convince you that if you are right about EN being non-ethical, then Karmo’s taxonomy has consequences that are undesirable for someone like you. So either way, whether you are right or whether I am right, Karmo’s taxonomy is incorrect.

²⁰ Here is the passage where Maitzen (1998, p. 361) motivates his taxonomy:

Our partitioning of ethical and non-ethical sentences should, among other things, capture the distinction between sentences that are merely *about* morality (and thus aren’t ethical sentences) and sentences that are genuinely *of*, or ‘belong to,’ morality (and thus are ethical sentences). [Consider the sentence] Morality often goes by the label ‘ethics.’ While, intuitively, [Torturing babies is wrong] is a sentence *of* morality [Morality often goes by the label ‘ethics’] is a sentence merely *about* morality. One way of capturing this distinction is to require that any genuinely ethical sentence entail that a given moral property (such as moral rightness or wrongness) is possessed by at least one object.

Karmo’s taxonomy can capture this distinction too. The truth of ‘Morality often goes by the label ‘ethics’’ does not vary with respect to different ethical standards. So Karmo does not classify this sentence as ethical. It seems to me that this is not enough of a motivation for a taxonomy that makes deriving an ‘Ought’ from an ‘Is’ so easy.

²¹ Huemer (p. 82), for example, says that EN “seems to be an evaluative statement.”

²² Thanks to Pete Graham for pressing me on this.

Perhaps you think EN is non-ethical because it is a sentence about sentences, that such sentences are about sentences exclusively, and are therefore non-ethical. If so, then consider the sentence ‘The sentence ‘All New Zealanders ought to be shot’ is false’. Call it *P*. If you think EN is non-ethical, then it seems to me that you should think *P* is non-ethical as well. EN says that every ethical sentence is false. *P* says that one particular ethical sentence is false. It seems to me that they should be classified together. But Karmo’s taxonomy classifies *P* as ethical. So if you are right about *P*, then Karmo’s taxonomy is incorrect. Now maybe you think that *P* is ethical while EN is not. If so, then your challenge is to find a feature EN possesses and *P* lacks that makes EN a non-ethical sentence.

Karmo’s Taxonomy Classifies all Necessary Sentences as Non-ethical

A related problem is that Karmo’s taxonomy misclassifies necessary and impossible ethical sentences. For any necessary sentence you pick, that sentence is going to be true at all possible worlds and with respect to all ethical standards. If it is impossible, then it is going to be false at all possible worlds and with respect to all ethical standards. So any necessary or impossible sentence is going to get classified, by Karmo, as non-ethical. But there are necessary and impossible ethical sentences. Therefore, Karmo’s taxonomy is incorrect.

My defense of the premise that there are necessary and impossible ethical sentences is, in addition to an appeal to intuition, this: There exists a prominent tradition in the literature that takes there to be necessary and impossible ethical sentences. This tradition includes Prior, Searle, Geach, and Huemer²³.

²³ Prior (1976, p. 89) says that such sentences “belong to the logic of ethics, or ‘deontic logic’ as it is sometimes called, but not to ethics itself.” However, he adds, “It will be appreciated that in piling on these conditions I am making my own job of deducing ethical conclusions a harder one.” Huemer (2007, p. 81) agrees that an interesting ‘Is’–‘Ought’ derivation will not have as its conclusion a necessary truth but still classifies such sentences as ethical. He says:

[T]he conclusion should not be a mere tautology such as ‘Right acts ought to be performed’ or ‘Bad things are bad’ (these statements trivially follow from any set of premises). I assume that an evaluative statement is non-trivial if there is at least one value system that rejects it.”

Geach (1977, p. 476) employs an allegedly necessary truth as a premise in his proposal. He classifies it as an ethical sentence but argues that although the sentence “is indeed an ought thesis,... it is vacuous, it imposes no logical restriction on our views of what in particular we ought or ought not to do.” Of this Huemer (2007, p. 76–7) says:

If [the sentence] were non-evaluative, the counter-example would seem to be in good order, but in fact [it] is evaluative. Nevertheless, Geach observes that [it] is an analytic truth.... So at least Geach seems to have shown how to go from non-evaluative statements and/or tautologies to a non-trivial evaluative conclusion.

So both Geach and Huemer think there are necessarily true ethical sentences.

Searle anticipates the claim that one of the premises in his proposal is an ethical sentence. He (1964, p. 46) says “[My opponent will claim] the derivation rests on the principle that one ought to keep one’s promises and that is a moral principle, hence evaluative. I don’t know whether “one ought to keep one’s promises” is a “moral” principle, but whether or not it is, it is also tautological.” The thing to take from this is that Searle thinks there can be tautological ethical sentences. Here is the upshot of this brief survey: There is a prominent tradition in the literature according to which there exist necessary and impossible ethical sentences.

Huemer on the Immunity of Karmo's Taxonomy

I have been arguing that Karmo's taxonomy gets some important classifications wrong. Perhaps my efforts have been misguided. Consider what Huemer (2007, p. 82) has to say about the matter:

[I]t might be that Karmo's sense of 'evaluative statement' differs from established usage—certainly it differs from Maitzen's usage—but this would not show Karmo's argument to be unsound either. In general, there is no question of 'falsifying' Karmo's characterization of evaluative statements, as long as that characterization is understood—as I believe it should be—as a stipulative definition.

Understand Karmo as picking out a certain set of sentences and naming them 'ethical'. If you are uneasy about calling these sentences 'ethical', then call them 'ethical*' instead. It might be that the classificatory conditions for genuinely ethical sentences are not identical to the classificatory conditions for ethical* sentences. But Karmo's taxonomy is stipulative. The aim is to show that ethical* sentences cannot be derived from non-ethical* ones. For this reason, simply demonstrating that the classificatory conditions for the two sorts of sentences differ does not threaten Karmo.

Huemer (2007, p. 82–3) anticipates the complaint that this makes Karmo's proof irrelevant. Here is what he says:

Given that the definition is purely stipulative, one may legitimately wonder whether Karmo's proof is relevant to what we were initially interested in—in particular, does Karmo's version of the is–ought gap help to establish that there is a problem for moral knowledge? I think it does. Karmo shows that there is no way of knowing evaluative statements by deduction from non-evaluative statements, in his sense of 'evaluative statements'. But Karmo's sense includes all the statements traditionally considered 'evaluative' by moral philosophers; thus, if it is difficult to know evaluative statements in Karmo's sense, it is also difficult to know evaluative statements in the traditional sense.

Notice in particular that Huemer says Karmo's "sense includes all the statements traditionally considered 'evaluative' by moral philosophers". It seems to me that Huemer is appealing to the following principle:

(H) If a taxonomy, *T*, is such that (i) there does not exist a sound 'Is'–'Ought' derivation with respect to *T*'s definition of ethical and (ii) any sentence that is genuinely ethical is also ethical with respect to *T*'s definition of ethical, then there does not exist a sound 'Is'–'Ought' derivation with respect to genuinely ethical sentences.

Karmo's taxonomy does not satisfy condition (ii). As I tried to show above, EN and all necessary and impossible ethical sentences are classified as non-ethical* by Karmo. Furthermore, necessary and impossible ethical sentences are "statements traditionally considered 'evaluative' by moral philosophers". Huemer himself is one of these philosophers. He also thinks that EN is an ethical

sentence²⁴. So you have sentences that are ethical but are not ethical with respect to Karmo's taxonomy. Condition (ii) is not satisfied²⁵. It seems to me that Huemer's defense of Karmo is unsuccessful²⁶.

²⁴ See footnote 21.

²⁵ A second problem with appealing to principle (H) is that it is false. (H) is a conditional: If (i) and (ii) are satisfied, then there are no sound 'Is'-'Ought' derivations with respect to really ethical sentences. We can test (H), therefore, by assuming that (H)'s consequent is false and asking whether it follows that (H)'s antecedent is false as well. So assume there are no sound 'Is'-'Ought' derivations. Does it follow that (i) and (ii) are not satisfied? No. I can show this by showing there is some way for (H)'s consequent to be false even though (i) and (ii) are both true. So, by hypothesis, I get to stipulate that there is a sound 'Is'-'Ought' derivation. If I can show there is some scenario according to which this hypothesis does not entail that (i) and (ii) are satisfied, then (H) is false.

So here is my strategy. I am going to state a taxonomy. I am going to come up with a scenario in which (H)'s consequent is false but (H)'s antecedent is true. This will show that principle (H) is false. Here is the taxonomy I propose:

A sentence, *S*, is ethical** if and only if either (a) *S* is ethical or (b) *S* is a premise in a sound argument that has a conclusion that is ethical.

The things Huemer says to defend Karmo's taxonomy can be used to defend the crazy taxonomy I have proposed. You cannot complain about the fact that the classificatory conditions of really ethical sentences differ from the classificatory conditions of ethical** sentences because my taxonomy is only stipulative. I am just picking out a certain set of sentences and I'm calling them 'ethical**'. I never claimed that these were one and the same as real ethical sentences. If this kind of reasoning gets Karmo off the hook then it gets me off the hook too.

Now, by hypothesis, there is an argument that is such that its conclusion is ethical and its premises are non-ethical. In order to show that (H) is false, I need to describe a scenario in which there is a sound 'Is'-'Ought' derivation with respect to real ethical sentences but there are no sound 'Is'-'Ought' derivations with respect to ethical** sentences and all ethical sentences are ethical** sentences. In other words I need to show that there is a scenario consistent with our hypothesis in which (i) and (ii) obtain. Here is one such scenario: Suppose there exists a sound argument, *A*, that is such that *A* has as its conclusion an ethical sentence and all of *A*'s premises are ethical**. Now, in order to be counted as ethical** a sentence must either (a) be ethical or (b) be a premise in a sound argument with an ethical conclusion. Suppose that all of *A*'s premises fail to satisfy (a). So they are not ethical. They only count as ethical** because they satisfy (b). In other words, they are ethical** because they are premises in a sound argument with an ethical conclusion. But the premises are not really ethical. So I have described a situation in which the consequent of (H) is false that is consistent with my crazy taxonomy. Nothing inherent to the taxonomy I have proposed requires that there are sound 'Is'-'Ought' derivations with respect to ethical sentences.

Now I need to show that even though (H)'s consequent is false on this scenario, its antecedent is true. You can show that there are not any sound derivations of ethical** sentences from non-ethical** ones. Here is how: Suppose there is a sound argument, *A*, that is such that all of *A*'s premises, *P*1 and *P*2 and..., are non-ethical** and *A*'s conclusion, *C*, is ethical**. By the definition of ethical**, either *C* is ethical or *C* is a premise in a sound argument that has an ethical sentence as its conclusion. Suppose that *C* is ethical. By the definition of ethical**, *P*1 and *P*2 and... are all ethical**. Suppose instead that *C* is a premise in a sound argument that has an ethical sentence as its conclusion. Then there is a sound argument that has *P*1 and *P*2 and... as premises and an ethical sentence as its conclusion. Namely, the argument that derives *C* from *P*1 and *P*2 and... and then goes on to derive an ethical sentence. So, by the definition of ethical**, *P*1 and *P*2 and... are all ethical**. Therefore, there is no sound argument, *A*, that is such that all of *A*'s premises, *P*1 and *P*2 and..., are non-ethical** and *A*'s conclusion, *C*, is ethical**. So condition (i) of (H) is satisfied.

You can also show that any sentence that is ethical is also ethical**. After all, any ethical sentence that you pick will satisfy condition (a) of my crazy taxonomy. What is more, satisfying either (a) or (b) is sufficient for a sentence to be classified as ethical**. So all ethical sentences are also ethical**. Condition (ii) of (H) is satisfied. The antecedent of (H) is true. So we have a taxonomy with respect to which (H)'s consequent, there are no sound 'Is'-'Ought' derivations is false but (H)'s antecedent, (i) and (ii) are satisfied, is true. Therefore principle (H) is false.

²⁶ A reviewer points out that the sort of 'Is'-'Ought' gap that Huemer is interested in is different than the traditional one—the one that I am interested in. Huemer's project is an epistemological one. He aims to show that there is no way to get from foundational empirical knowledge to knowledge of moral claims. We need intuition to account for moral knowledge. If this is right, then perhaps my argument does not bear on Huemer's project. In that case, I would say the following: It is still interesting to see if Huemer's point in defense of Karmo can be successfully made in the context of the debate that I am interested in—the traditional 'Is'-'Ought' gap.

Conclusion

Where does this discussion leave us? The most plausible ‘Is’–‘Ought’ derivations that have been proposed have a premise whose status is controversial. A taxonomy would help us to classify these sentences and to evaluate the derivations of which they are a part. Unfortunately, none of the current proposals are adequate. We still need a taxonomy of ethical sentences. There must be some way to tell an ‘Is’ from an ‘Ought’. But I do not know how to do it²⁷.

References

- Geach, P. T. (1977). Again the logic of “Ought”. *Philosophy*, 52, 473–476.
- Huemer, M. (2007). *Ethical intuitionism*. Palgrave Macmillan.
- Humberstone, I. L. (1982). First steps in philosophical taxonomy. *Canadian Journal of Philosophy*, 12, 476–478.
- Humberstone, I. L. (1996). A study in philosophical taxonomy. *Philosophical Studies*, 83, 121–169.
- Hume, D. [1739] (1992). *Treatise of human nature*. Reprint, Buffalo, N.Y.: Prometheus.
- Karmo, T. (1988). Some valid (but no sound) arguments trivially span the “Is”–“Ought” Gap’. *Mind*, 97, 252–257.
- Maitzen, S. (1998). ‘Closing the “Is”–“Ought” Gap’. *Canadian Journal of Philosophy*, 28, 349–366.
- Maitzen, S. (2006). The impossibility of local skepticism. *Philosophia*, 34, 453–464.
- Nelson, M. T. (2007). More bad news for the logical autonomy of ethics. *Canadian Journal of Philosophy*, 37, 203–216.
- Pigden, C. (1989). Logic and the autonomy of ethics. *Australasian Journal of Philosophy*, 67, 127–151.
- Pigden, C. (2007). Nihilism, nietzsche and the doppelganger problem. *Ethical Theory and Moral Practice*, 10, 441–456.
- Prior, A. N. (1976). The autonomy of ethics. In P. T. Geach, & A. J. P. Kenny (Eds.) *Papers in Logic and Ethics* (pp. 88–96). London: Duckworth.
- Searle, J. R. (1964). How to derive “Ought” from “Is”. *Philosophical Review*, 73, 43–58.
- Schurz, G. (1994). Hume’s is–ought thesis in logics with alethic–deontic bridge principles. *Logique et Analyse*, 37, 265–293.
- Schurz, G. (1997). *The Is–Ought problem: An investigation in philosophical logic*. Springer.

²⁷ I am indebted to Fred Feldman for uncommonly generous comments and guidance on numerous drafts of this paper. I also thank Pete Graham for helpful criticism. In addition I thank Gary Matthews, Lynne Baker, Phil Bricker, and a referee for this journal for comments that improved this paper.